BUSINESS CASE · NATURAL LANGUAGE PROCESSING

Twitter Named Entity Recognition BiLSTM-CRF & BERT

Recognising people, places, companies, products and more in noisy tweets — so trends can be understood without depending on hashtags.

Zendesk Security Breach Affects Twitter, Tumblr and Pinterest · Snapchat CEO talks about breach on Today Show · Made it back home to GA

company · person · tvshow · geo-loc — 10 fine-grained entity types, WNUT-16 Twitter corpus

0.437	0.395	+4.2 pts	17
BERT TEST ENTITY-F1	BILSTM-CRF TEST ENTITY-F1	GAIN FROM TRANSFORMER	TRAINING RUNS COMPARED

Guntuku Hariprasad
Scaler Academy · Business Case Submission

26 October 2026
TensorFlow 2.15 · Hugging Face Transformers

Contents

1. **Executive summary** — problem, approach, results, recommendations
2. **How this report maps to the evaluation criteria**
3. **Results at a glance** — all experiments in one table
4. **Part A — Data understanding, preprocessing & BiLSTM-CRF** (executed notebook)
    1. Loading the CoNLL data · 2. Structure analysis & EDA · 3. Preprocessing & TF Tokenizer · 4. Word2Vec embeddings · 5. BiLSTM + CRF model · 6. Experiments, hyper-parameters & early stopping · 7. Evaluation, alignment & predictions · 8. Saving
5. **Part B — Fine-tuning bert-base-uncased** (executed notebook)
    1. Data · 2. Model, config & tokenizer · 3. Tokenization & label alignment · 4. Joining sub-tokens · 5. Training, hyper-parameters & early stopping · 6. Test evaluation & comparison · 7. save_pretrained · 8. Own sentences · 9. Error analysis
6. **Answers to the business-case questions (Q1 – Q9)**
7. **Conclusion, recommendations & reproducibility**

1. Executive summary

Business problem. Twitter wants to understand trends and topics directly from tweet text, without depending on hashtags that are often missing, misspelled or misleading. The core capability is **Named Entity Recognition (NER)**: finding every mention of a *person, geo-location, company, facility, product, music artist, movie, sports team, TV show* or *other* entity in a tweet, with exact boundaries.

Data. WNUT-16 Twitter NER corpus in CoNLL/BIO format: 2,394 training tweets (46k tokens, 1,496 entities) and 3,849 test tweets (62k tokens, 3,473 entities), 10 types → 21 BIO labels. EDA shows what makes the task hard: **90–95% of tokens are `O`**, **85% of test entities never appear in training**, the entity mix shifts between train and test, and the text is noisy (URLs, @mentions, hashtags, slang, emoji).

Approach.

	Model A — BiLSTM-CRF	Model B — BERT
Input	TF `Tokenizer` (words, normalised) + char-level tokenizer	WordPiece `bert-base-uncased` tokenizer, labels aligned to first sub-token
Representation	Word2Vec (trained) / GloVe-Twitter (2B tweets) + char-CNN	12-layer pre-trained transformer, fine-tuned end-to-end
Output layer	Custom TensorFlow **CRF** with Viterbi decoding	Linear token-classification head + sub-token → word joining
Tuning	12 runs: direction, CRF, embeddings, char features, units, dropout, LR, batch	4 runs: learning rate, batch size, AdamW vs Adam, epochs; early stopping
Best test F1	**0.395** (P 0.433 / R 0.362)	**0.437** (P 0.432 / R 0.442)

Key findings.

1. **Pre-trained knowledge matters most.** For the BiLSTM-CRF, swapping embeddings trained on our 2k tweets for GloVe-Twitter embeddings took test F1 from 0.236 to 0.412, more than every architectural change combined. Fine-tuned BERT adds another **+4.2 F1 points** on top of that.
2. **Architecture still helps.** Bidirectional context (+3 pts), a CRF output layer (+2 pts) and character features (+5 pts) each improved the LSTM, and the CRF guarantees valid BIO sequences.
3. **Tokenization is the critical engineering step for BERT.** WordPiece removes the out-of-vocabulary problem, but it splits words into sub-tokens. Labels must be re-aligned for training, and predictions joined back to words for evaluation. Both steps are verified with assertions on every tweet.

4. **Early stopping** stopped the BiLSTM-CRF 12 epochs early, after its validation loss started rising, and saved about 30% of training time. With a 240-tweet validation set and a shifted test distribution, its effect on test F1 was small. For BERT, validation F1 kept improving through all 6 epochs, so it never triggered. See Q6.
5. **Where models fail:** rare and open-ended types (`movie`, `tvshow`, `product`, `other`) and semantically ambiguous names (`musicartist` ↔ `person`, `sportsteam` ↔ `geo-loc`).

Recommendation. Deploy the **fine-tuned transformer** as the entity extractor, and feed entity counts per time window into the trend-detection pipeline in place of hashtags. Next steps that will move the needle most: a Twitter-specific cased model (BERTweet / Twitter-RoBERTa), more labelled data for rare types (active learning), and a gazetteer / entity-linking layer for new names.

2. How this report maps to the evaluation criteria

Criterion (points)	Where it is addressed
Problem statement, importing the data & data-structure analysis (20)	Exec. summary; Q1; Part A §1 (defensive CoNLL parser, token tables), §2 (BIO integrity check, entity distributions, train/test shift, tweet & entity lengths, Twitter phenomena, OOV analysis); Q2
Correct model and tokenizer (10)	Part A §3 (TensorFlow `Tokenizer`, word + char), §4 (Word2Vec init); Part B §2 (`AutoConfig`, `AutoTokenizer`, `TFAutoModelForTokenClassification` for `bert-base-uncased`)
Data preprocessing & tokenization (30)	Part A §3 (normalisation, stratified train/val split, encoding, padding); Part B §3 (WordPiece analysis, [CLS]/[SEP], URL/@user decision, first-sub-token label alignment with assertions, dynamic padding & bucketing); Q4
Train the model (10)	Part A §5 (BiLSTM + custom TensorFlow CRF verified against brute force); Part B §5 (AdamW + warm-up fine-tuning)
Different hyper-parameters (10)	Part A §6 (12 experiments: direction, CRF, embeddings, char features, units, dropout, LR, batch, fine-tuning) & §6.1 early stopping; Part B §5 (LR, batch size, optimiser, epochs) & §5.1 early stopping; Q6
Make predictions — join sub-tokens; own sentences (20)	Part A §7 (alignment check, report, confusion matrix, own sentences); Part B §4 (sub-token → word joining + BIO repair), §6–§9 (test report, reload after `save_pretrained`, own sentences, error analysis)
Questions	Section 6 — Q1 to Q9

3. Results at a glance

All scores are **entity-level micro precision / recall / F1** (seqeval, exact span + type match) on the untouched WNUT-16 test set (3,849 tweets). Model selection always uses the **validation** F1 (10% of the training file), never the test set.

0.437
BERT TEST F1

0.432 / 0.442
BERT PRECISION / RECALL

0.395
BILSTM-CRF TEST F1

0.433 / 0.362
BILSTM-CRF PRECISION / RECALL

Part A — BiLSTM experiments (E1–E12)

experiment	epochs_run	val_F1	test_P	test_R	test_F1
E1 LSTM + softmax (w2v)	23	0.280	0.117	0.164	0.137
E2 BiLSTM + softmax (w2v)	40	0.399	0.137	0.210	0.166
E3 BiLSTM + CRF (random emb)	24	0.308	0.097	0.077	0.085
E4 BiLSTM + CRF (w2v)	35	0.436	0.183	0.193	0.188

experiment	epochs_run	val_F1	test_P	test_R	test_F1
E5 BiLSTM + CRF (w2v) + char-CNN	40	0.458	0.213	0.265	0.236
E6 BiLSTM + CRF (GloVe-Tw) + char-CNN	37	0.589	0.479	0.361	0.412
E7 E6 + units=64	33	0.545	0.486	0.348	0.405
E8 E6 + units=256	20	0.528	0.431	0.359	0.392
E9 E6 + dropout=0.3	27	0.571	0.469	0.363	0.410
E10 E6 + lr=3e-3	24	0.595	0.472	0.367	0.413
E11 E6 + batch=16	19	0.556	0.495	0.336	0.400
E12 E6 + fine-tuned embeddings ★ selected	**28**	**0.595**	**0.433**	**0.362**	**0.395**

Part B — BERT fine-tuning experiments (B1–B4)

experiment	epochs_run	val_F1	test_P	test_R	test_F1
B1 AdamW lr=5e-5 bs=16 ★ selected	**6**	**0.584**	**0.432**	**0.442**	**0.437**
B2 AdamW lr=3e-5 bs=16	5	0.510	0.335	0.374	0.353
B3 AdamW lr=5e-5 bs=32	5	0.404	0.266	0.314	0.288
B4 Adam (no decay/warm-up) lr=3e-5 bs=16	4	0.575	0.440	0.347	0.388

Part A — Data Understanding, Preprocessing & BiLSTM-CRF

Goal. Build a Named Entity Recognition (NER) system that tags tweets with 10 fine-grained entity types (`person, geo-loc, company, facility, product, musicartist, movie, sportsteam, tvshow, other`) so that Twitter can surface trends and topics *without relying on user hashtags*.

This notebook covers: data loading & structure analysis, exploratory data analysis, preprocessing, a TensorFlow tokenizer, Word2Vec-initialised embeddings, and a **BiLSTM + CRF** sequence tagger with hyper-parameter experiments. Part B (separate notebook) fine-tunes `bert-base-uncased`.

In [1]
```
import os, re, json, time, random, warnings
from collections import Counter, defaultdict

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

os.environ["TF_CPP_MIN_LOG_LEVEL"] = "3"
warnings.filterwarnings("ignore")

import tensorflow as tf
tf.get_logger().setLevel("ERROR")
from gensim.models import Word2Vec
from seqeval.metrics import classification_report, f1_score, precision_score, recall_score
from sklearn.model_selection import train_test_split

SEED = 42
tf.keras.utils.set_random_seed(SEED)
random.seed(SEED); np.random.seed(SEED)

sns.set_theme(style="whitegrid", context="notebook")
PALETTE = ["#1D9BF0", "#0F1419", "#F91880", "#00BA7C", "#FFD400", "#7856FF", "#FF7A00", "#536471",
"#8ECDF8", "#C1A3FF"]
plt.rcParams.update({"figure.dpi": 110, "axes.titleweight": "bold", "axes.titlesize": 12})
pd.set_option("display.max_colwidth", 110); pd.set_option("display.width", 160);
pd.set_option("display.max_columns", 30)

DATA_DIR, OUT_DIR = "../data", "../outputs"
os.makedirs(OUT_DIR, exist_ok=True)
print("TensorFlow", tf.__version__, "| GPUs:", tf.config.list_physical_devices("GPU"))
```

```
TensorFlow 2.15.0 | GPUs: []
```

1. Loading the data (CoNLL format)

Each line holds `token<TAB>label`; tweets are separated by a blank line. Labels follow the **BIO** scheme (`B-` = beginning of an entity, `I-` = inside/continuation, `O` = outside any entity). The parser below is defensive: it handles Windows line endings, stray whitespace and validates every line.

In [2]
```
def read_conll(path):
    """Parse a CoNLL file into parallel lists of token sequences and tag sequences."""
    sentences, tags, toks, labs, bad = [], [], [], [], 0
    with open(path, encoding="utf-8") as f:
        for line in f:
            line = line.rstrip("\r\n")
            if not line.strip():                    # blank line -> sentence boundary
                if toks:
```

```
                    sentences.append(toks); tags.append(labs); toks, labs = [], []
                continue
            parts = line.split("\t") if "\t" in line else line.split()
            if len(parts) < 2:
                bad += 1; continue
            toks.append(parts[0]); labs.append(parts[-1])   # first column = word, last column = label
    if toks:
        sentences.append(toks); tags.append(labs)
    return sentences, tags, bad

train_sents, train_tags, bad_tr = read_conll(f"{DATA_DIR}/wnut 16.txt.conll")
test_sents, test_tags, bad_te = read_conll(f"{DATA_DIR}/wnut 16test.txt.conll")
print(f"Train file : {len(train_sents):>5} tweets | {sum(map(len, train_sents)):>6} tokens | malformed
lines: {bad_tr}")
print(f"Test file  : {len(test_sents):>5} tweets | {sum(map(len, test_sents)):>6} tokens | malformed lines:
{bad_te}")
```

```
Train file :  2394 tweets |  46469 tokens | malformed lines: 0
Test file  :  3850 tweets |  61908 tokens | malformed lines: 0
```

In [3]

```
# Token / label view of two short training tweets (one row per CoNLL line, shown transposed)
short = [i for i, (s, t) in enumerate(zip(train_sents, train_tags)) if len(s) <= 13 and sum(x != "O" for x
in t) >= 2][:2]
for i in short:
    display(pd.DataFrame({"token": train_sents[i], "label": train_tags[i]}).T)
```

	0	1	2	3	4	5	6	7	8	9	10
token	'	Breaking	Dawn	'	Returns	to	Vancouver	on	January	11th	http://bit.ly/dbDMs8
label	O	B-movie	I-movie	O	O	O	B-geo-loc	O	O	O	O

	0	1	2	3	4	5	6	7	8	9	10	11
token	#SIUC	Whats	the	Plan	For	tonight	??	Whos	Goin	to	the	BLAST
label	O	O	O	O	O	O	O	O	O	O	B-facility	I-facility

In [4]

```
def to_frame(sents, tags, split):
    return pd.DataFrame([(split, i, j, w, t) for i, (s, ts) in enumerate(zip(sents, tags))
                         for j, (w, t) in enumerate(zip(s, ts))],
                        columns=["split", "sent_id", "position", "token", "tag"])

df = pd.concat([to_frame(train_sents, train_tags, "train"), to_frame(test_sents, test_tags, "test")],
ignore_index=True)
df["entity_type"] = df.tag.str.replace(r"^[BI]-", "", regex=True)
df.info()
df.head()
```

```
<class 'pandas.DataFrame'>
RangeIndex: 108377 entries, 0 to 108376
Data columns (total 6 columns):
 #   Column       Non-Null Count   Dtype
---  ------       --------------   -----
 0   split        108377 non-null  str
 1   sent_id      108377 non-null  int64
 2   position     108377 non-null  int64
 3   token        108377 non-null  str
 4   tag          108377 non-null  str
 5   entity_type  108377 non-null  str
dtypes: int64(2), str(4)
memory usage: 5.0 MB
```

	split	sent_id	position	token	tag	entity_type
0	train	0	0	@SammieLynnsMom	O	O

	split	sent_id	position	token	tag	entity_type
1	train	0	1	@tg10781	O	O
2	train	0	2	they	O	O
3	train	0	3	will	O	O
4	train	0	4	be	O	O

2. Data structure analysis & EDA

2.1 Label inventory and BIO integrity check

In [5]

```
ENTITY_TYPES = sorted({t[2:] for t in df.tag if t != "O"})
LABELS = ["O"] + [f"{p}-{e}" for e in ENTITY_TYPES for p in ("B", "I")]
print(f"{len(ENTITY_TYPES)} entity types -> {len(LABELS)} BIO labels")
print(ENTITY_TYPES)

def bio_violations(tag_seqs):
    """Count I-X tags that are not preceded by B-X or I-X (invalid BIO transitions)."""
    v = 0
    for seq in tag_seqs:
        prev = "O"
        for t in seq:
            if t.startswith("I-") and prev[2:] != t[2:]:
                v += 1
            prev = t
    return v

print("BIO violations  train:", bio_violations(train_tags), "| test:", bio_violations(test_tags))
print("Null/empty tokens:", (df.token.str.strip() == "").sum(), "| duplicated tweets in train:",
      len(train_sents) - len({" ".join(s) for s in train_sents}))
```

```
10 entity types -> 21 BIO labels
['company', 'facility', 'geo-loc', 'movie', 'musicartist', 'other', 'person', 'product', 'sportsteam',
'tvshow']
BIO violations  train: 0 | test: 0
Null/empty tokens: 0 | duplicated tweets in train: 99
```

In [6]

```
def extract_entities(tokens, tags):
    """Return list of (entity_type, entity_text, length) spans from a BIO sequence."""
    ents, cur, typ = [], [], None
    for w, t in zip(tokens + [""], tags + ["O"]):     # sentinel "O" flushes the last open entity
        if t.startswith("B-") or t == "O" or (t.startswith("I-") and t[2:] != typ):
            if cur:
                ents.append((typ, " ".join(cur), len(cur)))
            cur, typ = ([w], t[2:]) if t != "O" else ([], None)
        else:
            cur.append(w)
    return ents

ent_rows = [(split, *e) for split, S, T in [("train", train_sents, train_tags), ("test", test_sents,
test_tags)]
            for s, t in zip(S, T) for e in extract_entities(s, t)]
ents = pd.DataFrame(ent_rows, columns=["split", "type", "text", "length"])

summary = pd.DataFrame({
    "tweets": [len(train_sents), len(test_sents)],
    "tokens": [sum(map(len, train_sents)), sum(map(len, test_sents))],
    "entities": [(ents.split == "train").sum(), (ents.split == "test").sum()],
    "tweets_with_entity": [sum(any(t != "O" for t in ts) for ts in train_tags),
                           sum(any(t != "O" for t in ts) for ts in test_tags)],
    "pct_O_tokens": [round(100 * np.mean([t == "O" for ts in train_tags for t in ts]), 2),
                     round(100 * np.mean([t == "O" for ts in test_tags for t in ts]), 2)],
```

```
}, index=["train", "test"])
summary
```

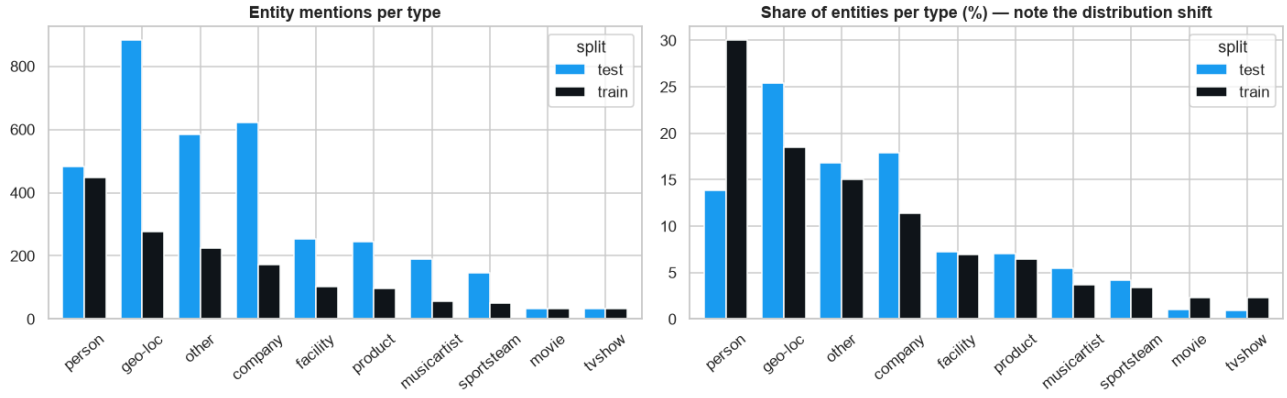
	tweets	tokens	entities	tweets_with_entity	pct_O_tokens
train	2394	46469	1496	918	94.70
test	3850	61908	3473	2014	90.38

2.2 Entity-type distribution (train vs test)

In [7]

```
ct = ents.groupby(["type", "split"]).size().unstack().fillna(0).astype(int)
ct = ct.loc[ct.train.sort_values(ascending=False).index]
pct = ct / ct.sum() * 100

fig, axes = plt.subplots(1, 2, figsize=(13, 4.2))
ct.plot.bar(ax=axes[0], color=[PALETTE[1], PALETTE[0]][::-1], width=0.75)
axes[0].set_title("Entity mentions per type"); axes[0].set_xlabel(""); axes[0].tick_params(axis="x",
rotation=40)
pct.plot.bar(ax=axes[1], color=[PALETTE[1], PALETTE[0]][::-1], width=0.75)
axes[1].set_title("Share of entities per type (%) — note the distribution shift"); axes[1].set_xlabel("")
axes[1].tick_params(axis="x", rotation=40)
plt.tight_layout(); plt.show()
pd.concat([ct, pct.round(1).add_suffix(" %")], axis=1)
```



split	test	train	test %	train %
type				
person	482	449	13.9	30.0
geo-loc	882	276	25.4	18.4
other	584	225	16.8	15.0
company	621	171	17.9	11.4
facility	253	104	7.3	7.0
product	246	97	7.1	6.5
musicartist	191	55	5.5	3.7
sportsteam	147	51	4.2	3.4
movie	34	34	1.0	2.3
tvshow	33	34	1.0	2.3

Observations

- The dataset is **heavily imbalanced**: ~94% of tokens are `O`. Accuracy is therefore a misleading metric — we use **entity-level precision / recall / F1 (seqeval, CoNLL-style exact span match)** throughout.
- There is a **distribution shift** between train and test: `company`, `geo-loc` and `product` are far more frequent in the test set, while `person` dominates train. Models must generalise beyond memorised names.

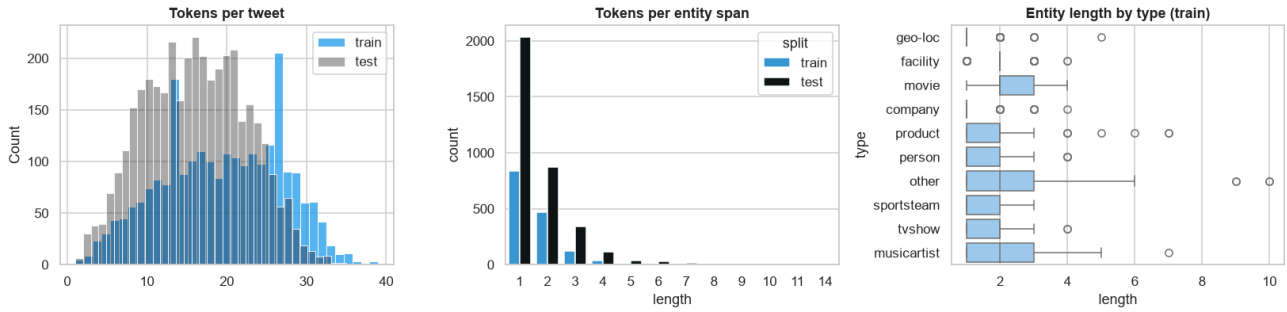
- Rare classes (`movie`, `tvshow`, `sportsteam`, `musicartist`) have < 60 training examples each — expected to be the hardest.

2.3 Tweet length, entity length and Twitter-specific phenomena

In [8]
```
fig, axes = plt.subplots(1, 3, figsize=(15, 3.8))
sns.histplot([len(s) for s in train_sents], bins=35, ax=axes[0], color=PALETTE[0], label="train")
sns.histplot([len(s) for s in test_sents], bins=35, ax=axes[0], color=PALETTE[1], alpha=.35, label="test")
axes[0].set_title("Tokens per tweet"); axes[0].legend()
sns.countplot(data=ents, x="length", hue="split", ax=axes[1], palette=[PALETTE[0], PALETTE[1]])
axes[1].set_title("Tokens per entity span")
sns.boxplot(data=ents[ents.split == "train"], y="type", x="length", ax=axes[2], color=PALETTE[8])
axes[2].set_title("Entity length by type (train)")
plt.tight_layout(); plt.show()

def token_kind(w):
    if re.match(r"https?://|www\.", w): return "URL"
    if w.startswith("@") and len(w) > 1: return "@mention"
    if w.startswith("#") and len(w) > 1: return "#hashtag"
    if not re.search(r"[A-Za-z0-9]", w): return "punct/emoji"
    if w[:1].isupper(): return "Capitalised"
    return "lowercase"

df["kind"] = df.token.map(token_kind)
kind_tab = pd.crosstab(df.kind, df.tag != "O", normalize="columns").mul(100).round(1)
kind_tab.columns = ["% of O tokens", "% of entity tokens"]
kind_tab.sort_values("% of entity tokens", ascending=False)
```



	% of O tokens	% of entity tokens
kind		
Capitalised	17.6	83.3
lowercase	59.9	11.2
#hashtag	2.1	3.9
punct/emoji	14.2	1.5
URL	2.8	0.1
@mention	3.4	0.0

In [9]
```
print("Max tweet length (tokens): train", max(map(len, train_sents)), "| test", max(map(len, test_sents)))
print("\nMost frequent entities per type (train):")
top = (ents[ents.split == "train"].groupby("type").text
       .agg(lambda s: ", ".join(f"{k} ({v})" for k, v in Counter(s.str.lower()).most_common(5))))
top.to_frame("top-5 surface forms")
```

```
Max tweet length (tokens): train 39 | test 35

Most frequent entities per type (train):
```

	top-5 surface forms
type	
company	twitter (32), facebook (13), youtube (13), ustream (5), playboy (4)
facility	visions lounge (3), disney world (2), costa lounge (2), lhs (2), frat house hattiesburg (2)
geo-loc	chicago (8), london (6), ga (5), nyc (5), dc (4)
movie	the town (3), gigli (2), winter 's bone (2), breaking dawn (1), piranha 3d (1)
musicartist	kings of leon (2), cocofunka (2), green day (2), 30stm (2), kings and queens (2)
other	yom kippur (6), christmas (4), constitution day (3), london fashion week (3), xmas (3)
person	justin bieber (16), pope (12), lindsay lohan (10), god (9), justin (6)
product	ipad (6), htc desire z (4), club penguin (3), iphone (3), pumpkin moonshine (2)
sportsteam	eagles (3), the lions (2), jv soccer (2), eyeopener (2), sharks (2)
tvshow	friday night lights (2), indiesent exposure (2), private practice (2), true blood (2), demos (1)

In [10]
```
train_vocab = Counter(w.lower() for s in train_sents for w in s)
test_tokens = [w.lower() for s in test_sents for w in s]
test_ent_tokens = [w.lower() for s, ts in zip(test_sents, test_tags) for w, t in zip(s, ts) if t != "O"]
oov_all = np.mean([w not in train_vocab for w in test_tokens]) * 100
oov_ent = np.mean([w not in train_vocab for w in test_ent_tokens]) * 100
seen_ents = set(ents[ents.split == "train"].text.str.lower())
unseen = np.mean([e not in seen_ents for e in ents[ents.split == "test"].text.str.lower()]) * 100
print(f"Train vocabulary size: {len(train_vocab):,}")
print(f"Test token OOV rate          : {oov_all:5.1f}%")
print(f"Test ENTITY-token OOV rate   : {oov_ent:5.1f}%")
print(f"Test entities never seen in train: {unseen:5.1f}%")
```

```
Train vocabulary size: 9,068
Test token OOV rate          :  30.1%
Test ENTITY-token OOV rate   :  52.1%
Test entities never seen in train:  85.1%
```

Key EDA insights that drive the modelling choices

1. **Short, noisy text** — tweets are ≤ 39 tokens, with URLs, @mentions, hashtags, emoji and slang. URLs and @mentions are (almost) never entities → we normalise them to placeholders.
2. **Capitalisation is a strong but imperfect cue** — 83% of entity tokens are capitalised vs only 18% of `O` tokens, yet 11% of entity tokens are lower-case (e.g. `pinterest`). → we lower-case the *word* input for better embedding coverage but add **character-level CNN features** on the raw token so casing and morphology are still visible.
3. **Very high OOV** — 30% of test tokens and 52% of test *entity* tokens never occur in training, and **85% of test entities were never seen in training**. Memorising names cannot work. → **pre-trained / sub-word representations** (Word2Vec, GloVe-Twitter, BERT WordPiece) are essential.
4. **Multi-token entities** and strict BIO constraints (`I-x` must follow `B-x` / `I-x`). → a **CRF** output layer that models label transitions jointly is preferable to independent softmax decisions.
5. **Severe class imbalance and train→test shift** (90–95% `O` tokens; `geo-loc` / `company` share grows in test) → evaluate with entity-level F1 (seqeval), not token accuracy, and expect test scores below validation scores.

3. Preprocessing for the LSTM model

1. **Normalisation** (word branch): URLs → `<url>`, @mentions → `<user>`, `#tag` → `tag`, digits → `0`, lower-case. Raw tokens are kept for the **character branch**, so casing / `#` / `@` cues are not lost.
2. **Train / validation split** — the provided training file is split 90/10 (stratified on "tweet contains an entity"); the provided test file is kept untouched for the final evaluation.
3. **TensorFlow** `Tokenizer` fitted on the training split only (no leakage) for words, and a second char-level tokenizer.

4. **Label encoding** to integer ids; **padding** to a fixed length of 40 (longest tweet = 39 tokens, so nothing is truncated).

In [11]
```
def normalise(tok):
    if re.match(r"https?://|www\.", tok, re.I): return "<url>"
    if tok.startswith("@") and len(tok) > 1:   return "<user>"
    if tok.startswith("#") and len(tok) > 1:   tok = tok[1:]
    return re.sub(r"\d", "0", tok.lower())

has_ent = [int(any(t != "O" for t in ts)) for ts in train_tags]
tr_idx, va_idx = train_test_split(np.arange(len(train_sents)), test_size=0.10, random_state=SEED,
stratify=has_ent)
X_tr_raw, y_tr_raw = [train_sents[i] for i in tr_idx], [train_tags[i] for i in tr_idx]
X_va_raw, y_va_raw = [train_sents[i] for i in va_idx], [train_tags[i] for i in va_idx]
X_te_raw, y_te_raw = test_sents, test_tags
print(f"train {len(X_tr_raw)} | validation {len(X_va_raw)} | test {len(X_te_raw)} tweets")
print("Example normalisation:", [(w, normalise(w)) for w in ["@Snapchat", "#Twitter", "http://bit.ly/x",
"2015", "Zendesk"]])
```

```
train 2154 | validation 240 | test 3850 tweets
Example normalisation: [('@Snapchat', '<user>'), ('#Twitter', 'twitter'), ('http://bit.ly/x', '<url>'),
('2015', '0000'), ('Zendesk', 'zendesk')]
```

In [12]
```
MAX_LEN, MAX_WORD_LEN = 40, 20

word_tok = tf.keras.preprocessing.text.Tokenizer(filters="", lower=False, oov_token="<oov>", split=" ")
word_tok.fit_on_texts([" ".join(normalise(w) for w in s) for s in X_tr_raw])
char_tok = tf.keras.preprocessing.text.Tokenizer(char_level=True, lower=False, oov_token="\x01")
char_tok.fit_on_texts([w for s in X_tr_raw for w in s])

tag2id = {t: i for i, t in enumerate(LABELS)}
id2tag = {i: t for t, i in tag2id.items()}
VOCAB, CHARS, N_TAGS = len(word_tok.word_index) + 1, len(char_tok.word_index) + 1, len(LABELS)
print(f"word vocab: {VOCAB:,} | char vocab: {CHARS} | tags: {N_TAGS}")

def encode(sents, tags=None):
    w = [word_tok.texts_to_sequences([" ".join(normalise(t) for t in s)])[0] for s in sents]
    Xw = tf.keras.preprocessing.sequence.pad_sequences(w, maxlen=MAX_LEN, padding="post",
truncating="post")
    Xc = np.zeros((len(sents), MAX_LEN, MAX_WORD_LEN), dtype="int32")
    for i, s in enumerate(sents):
        for j, tok in enumerate(s[:MAX_LEN]):
            ids = char_tok.texts_to_sequences([list(tok)])[0][:MAX_WORD_LEN]
            Xc[i, j, :len(ids)] = ids
    lens = np.array([min(len(s), MAX_LEN) for s in sents])
    Y = None
    if tags is not None:
        Y = tf.keras.preprocessing.sequence.pad_sequences([[tag2id[t] for t in ts] for ts in tags],
                                                          maxlen=MAX_LEN, padding="post",
value=tag2id["O"])
    return Xw, Xc, lens, Y

Xw_tr, Xc_tr, L_tr, Y_tr = encode(X_tr_raw, y_tr_raw)
Xw_va, Xc_va, L_va, Y_va = encode(X_va_raw, y_va_raw)
Xw_te, Xc_te, L_te, Y_te = encode(X_te_raw, y_te_raw)
print("word ids:", Xw_tr.shape, "| char ids:", Xc_tr.shape, "| labels:", Y_tr.shape)
i = next(k for k, t in enumerate(y_tr_raw) if 8 <= len(t) <= 12 and any(x != "O" for x in t))   # short
tweet with an entity
pd.DataFrame({"raw": X_tr_raw[i], "normalised": [normalise(w) for w in X_tr_raw[i]],
              "word_id": Xw_tr[i][:L_tr[i]], "label": y_tr_raw[i], "label_id": Y_tr[i][:L_tr[i]]}).T
```

```
word vocab: 6,780 | char vocab: 94 | tags: 21
```

```
word ids: (2154, 40) | char ids: (2154, 40, 20) | labels: (2154, 40)
```

	0	1	2	3	4	5	6	7
raw	Ceremony	Sunday	to	mark	Battle	of	Britain	http://dlvr.it/5RhwW
normalised	ceremony	sunday	to	mark	battle	of	britain	<url>
word_id	2565	109	7	676	790	20	2566	14
label	O	O	O	O	B-other	I-other	I-other	O
label_id	0	0	0	0	11	12	12	0

4. Word2Vec embeddings

We train a **skip-gram Word2Vec** model (gensim) on the normalised training tweets and use it to **initialise** the embedding layer. Because the training corpus is tiny (~2k tweets), we additionally compare against **GloVe-Twitter-100** vectors pre-trained on 2B tweets (loaded through the same gensim `KeyedVectors` interface) to quantify how much large-scale, in-domain pre-training helps.

In [13]
```
EMB_DIM = 100
w2v = Word2Vec(sentences=[[normalise(w) for w in s] for s in X_tr_raw], vector_size=EMB_DIM, window=5,
               min_count=1, sg=1, negative=10, epochs=50, seed=SEED, workers=1)
print("Word2Vec vocabulary:", len(w2v.wv))
for q in ["<user>", "twitter", "london", "game"]:
    if q in w2v.wv:
        print(f"  most similar to {q!r:10}:", [w for w, _ in w2v.wv.most_similar(q, topn=6)])

def build_matrix(kv, dim=EMB_DIM):
    rng = np.random.default_rng(SEED)
    M = rng.normal(0, 0.1, (VOCAB, dim)).astype("float32"); M[0] = 0
    hits = 0
    for word, idx in word_tok.word_index.items():
        key = {"<user>": "<user>", "<url>": "<url>"}.get(word, word)
        if key in kv:
            M[idx] = kv[key]; hits += 1
    return M, hits / (VOCAB - 1)

W2V_MATRIX, cov = build_matrix(w2v.wv)
print(f"Word2Vec embedding coverage of tokenizer vocabulary: {cov:.1%}")
```

```
Word2Vec vocabulary: 6778
  most similar to '<user>'  : ['rt', ':', 'grumpy', 'spirited', "/lebowski'd", 'knew']
  most similar to 'twitter' : ['spew', 'jackie', 'joins', 'trending', 'yahoo', 'beginners']
  most similar to 'london'  : ['clogging', 'england', 'nighter', 'twisted', 'hyde', 'travelling']
  most similar to 'game'    : ['tooooonight', 'foooooootball', 'wolves', 'havnt', 'hall', 'shs']
Word2Vec embedding coverage of tokenizer vocabulary: 100.0%
```

In [14]
```
import gensim.downloader as api
glove = api.load("glove-twitter-100")       # 1.2M-word vocabulary trained on 2B tweets
GLOVE_MATRIX, cov_g = build_matrix(glove)
test_vocab = {normalise(w) for s in X_te_raw for w in s}
print(f"GloVe-Twitter coverage: train vocab {cov_g:.1%} | test vocab {np.mean([w in glove for w in
test_vocab]):.1%} "
      f"(vs Word2Vec-on-train test coverage {np.mean([w in w2v.wv for w in test_vocab]):.1%})")
```

```
GloVe-Twitter coverage: train vocab 85.4% | test vocab 81.1% (vs Word2Vec-on-train test coverage 27.5%)
```

Because the **embedding matrix is indexed by the training tokenizer**, test words outside that vocabulary would map to `<oov>` even if GloVe knows them. To exploit pre-trained vectors fully, the GloVe variant uses an **extended vocabulary**: words from the pre-trained table that appear in the *unlabelled* text of the evaluation data are appended to the embedding matrix and **kept frozen**, which is a standard practice (no labels are used). For the Word2Vec experiments we keep the strict train-only vocabulary.

In [15]
```
def extend_vocab_with(kv):
    """Return a new tokenizer index + matrix covering train vocab ∪ (eval words found in kv)."""
    index = dict(word_tok.word_index)
    extra = sorted({normalise(w) for s in X_va_raw + X_te_raw for w in s} - set(index))
    extra = [w for w in extra if w in kv]
    for w in extra:
        index[w] = len(index) + 1
    rng = np.random.default_rng(SEED)
    M = rng.normal(0, 0.1, (len(index) + 1, EMB_DIM)).astype("float32"); M[0] = 0
    for w, i in index.items():
        if w in kv: M[i] = kv[w]
    return index, M, len(extra)

EXT_INDEX, EXT_MATRIX, n_extra = extend_vocab_with(glove)
print(f"Extended vocabulary: +{n_extra:,} pre-trained words -> {len(EXT_INDEX)+1:,} rows")

def encode_words_ext(sents):
    oov = EXT_INDEX["<oov>"]
    seqs = [[EXT_INDEX.get(normalise(w), oov) for w in s] for s in sents]
    return tf.keras.preprocessing.sequence.pad_sequences(seqs, maxlen=MAX_LEN, padding="post")

Xg_tr, Xg_va, Xg_te = encode_words_ext(X_tr_raw), encode_words_ext(X_va_raw), encode_words_ext(X_te_raw)
print("OOV rate on test words: train-vocab tokenizer "
      f"{np.mean(Xw_te[Xw_te>0]==1):.1%} -> extended {np.mean(Xg_te[Xg_te>0]==EXT_INDEX['<oov>']):.1%}")
```

```
Extended vocabulary: +6,136 pre-trained words -> 12,916 rows
```

```
OOV rate on test words: train-vocab tokenizer 21.0% -> extended 3.6%
```

5. Model — BiLSTM + CRF

```
word ids ──► Embedding (Word2Vec / GloVe init) ─┐
                                                ├─► concat ─► Dropout ─► BiLSTM ─► Dense(K) = emission scores ─► CRF
char ids ──► Char-Embedding ─► Conv1D ─► MaxPool ┘                                                              (+ K×K transition matrix)
```

- **Bidirectional LSTM** — each token representation sees both left *and* right context (e.g. "*at* **Disney** *world*"). A unidirectional model only sees the past, which hurts the `B-` / `I-` decision on the first token of an entity.
- **CRF layer** — instead of independent softmax decisions it learns a **transition matrix** and maximises the likelihood of the *whole* tag sequence; decoding uses the **Viterbi** algorithm, which also prevents invalid sequences like `O → I-person`. The CRF is implemented below directly in TensorFlow (forward algorithm for the partition function + Viterbi decoding), so it has no dependency on the deprecated `tensorflow-addons` package.

In [16]
```
class CRF(tf.keras.layers.Layer):
    """Linear-chain CRF: learns transition, start and end scores; computes NLL and Viterbi decoding."""
    def __init__(self, num_tags, **kw):
        super().__init__(**kw); self.num_tags = num_tags

    def build(self, input_shape):
        init = tf.keras.initializers.RandomUniform(-0.1, 0.1)
        self.trans = self.add_weight(name="transitions", shape=(self.num_tags, self.num_tags),
initializer=init)
        self.start = self.add_weight(name="start", shape=(self.num_tags,), initializer=init)
        self.end = self.add_weight(name="end", shape=(self.num_tags,), initializer=init)

    def nll(self, emissions, tags, lengths):
        """Negative log-likelihood of gold tag sequences. emissions [B,T,K], tags [B,T], lengths [B]."""
        T = tf.shape(emissions)[1]
        mask = tf.sequence_mask(lengths, T, dtype=emissions.dtype)                    # [B,T]
        tags = tf.cast(tags, tf.int32)
```

```
        # --- score of the gold path
        unary = tf.reduce_sum(tf.gather(emissions, tags, batch_dims=2) * mask, axis=1)
        pair = tf.gather_nd(self.trans, tf.stack([tags[:, :-1], tags[:, 1:]], axis=-1))   # [B,T-1]
        binary = tf.reduce_sum(pair * mask[:, 1:], axis=1)
        last = tf.gather(tags, lengths - 1, batch_dims=1)
        gold = unary + binary + tf.gather(self.start, tags[:, 0]) + tf.gather(self.end, last)
        # --- log partition function (forward algorithm)
        alpha = self.start[None] + emissions[:, 0]
        em_t = tf.transpose(emissions, [1, 0, 2])[1:]; m_t = tf.transpose(mask)[1:]
        def step(a, x):
            e, m = x
            nxt = tf.reduce_logsumexp(a[:, :, None] + self.trans[None] + e[:, None, :], axis=1)
            return m[:, None] * nxt + (1 - m[:, None]) * a
        alpha = tf.foldl(step, (em_t, m_t), initializer=alpha)
        log_z = tf.reduce_logsumexp(alpha + self.end[None], axis=1)
        return log_z - gold

    def viterbi(self, emissions, lengths):
        """Numpy Viterbi decoding -> list of best tag-id paths (unpadded)."""
        A, s, e = self.trans.numpy(), self.start.numpy(), self.end.numpy()
        paths = []
        for em, n in zip(emissions, lengths):
            score, back = s + em[0], []
            for t in range(1, n):
                cand = score[:, None] + A
                back.append(cand.argmax(0)); score = cand.max(0) + em[t]
            best = [int((score + e).argmax())]
            for bp in reversed(back):
                best.append(int(bp[best[-1]]))
            paths.append(best[::-1])
        return paths


class CRFTagger(tf.keras.Model):
    """Wraps an emission network with a CRF; custom train/test steps optimise the sequence NLL."""
    def __init__(self, encoder, num_tags):
        super().__init__(); self.encoder, self.crf = encoder, CRF(num_tags)
        self.crf.build(None); self.loss_tracker = tf.keras.metrics.Mean(name="loss")

    def call(self, inputs, training=False):
        return self.encoder(inputs, training=training)

    @property
    def metrics(self): return [self.loss_tracker]

    def _loss(self, x, y, training):
        inputs, lengths = x
        em = self(inputs, training=training)
        return tf.reduce_mean(self.crf.nll(em, y, lengths))

    def train_step(self, data):
        x, y = data
        with tf.GradientTape() as tape:
            loss = self._loss(x, y, True)
        grads = tape.gradient(loss, self.trainable_variables)
        self.optimizer.apply_gradients(zip(grads, self.trainable_variables))
        self.loss_tracker.update_state(loss); return {"loss": self.loss_tracker.result()}

    def test_step(self, data):
        x, y = data
        self.loss_tracker.update_state(self._loss(x, y, False)); return {"loss":
self.loss_tracker.result()}

    def decode(self, inputs, lengths, batch_size=256):
        em = self.encoder.predict(inputs, batch_size=batch_size, verbose=0)
        return self.crf.viterbi(em, lengths)
```

In [17]
```
# Sanity check of the CRF: brute-force enumeration of all paths must give the same log-partition & Viterbi path
crf = CRF(3); crf.build(None)
em = tf.random.normal((1, 4, 3)); gold = tf.constant([[0, 1, 2, 1]])
import itertools
scores = {p: float(crf.start[p[0]] + sum(em[0, t, p[t]] for t in range(4)) +
                   sum(crf.trans[p[t-1], p[t]] for t in range(1, 4)) + crf.end[p[-1]])
          for p in itertools.product(range(3), repeat=4)}
log_z = np.log(np.sum(np.exp(list(scores.values()))))
print("NLL  (CRF layer) :", round(float(crf.nll(em, gold, tf.constant([4]))[0]), 5))
print("NLL  (brute force):", round(log_z - scores[(0, 1, 2, 1)], 5))
print("Viterbi path:", crf.viterbi(em.numpy(), [4])[0], "| brute-force argmax:", list(max(scores, key=scores.get)))
```

```
NLL  (CRF layer) : 5.58392
NLL  (brute force): 5.58392
Viterbi path: [0, 2, 2, 0] | brute-force argmax: [0, 2, 2, 0]
```

5.1 Model factory, entity-F1 callback and training helper

Early stopping monitors **validation entity-level F1** (the metric we actually care about), computed with `seqeval` after Viterbi decoding at the end of each epoch; the best weights are restored at the end.

In [18]
```
def build_model(cfg):
    """Build a (Bi)LSTM tagger with optional char-CNN features and a CRF or softmax output layer."""
    emb_kind = cfg.get("emb", "w2v")
    vocab = len(EXT_INDEX) + 1 if emb_kind == "glove" else VOCAB
    matrix = {"w2v": W2V_MATRIX, "glove": EXT_MATRIX, "random": None}[emb_kind]
    w_in = tf.keras.Input((MAX_LEN,), dtype="int32", name="words")
    init = tf.keras.initializers.Constant(matrix) if matrix is not None else "uniform"
    emb = tf.keras.layers.Embedding(vocab, EMB_DIM, mask_zero=True, embeddings_initializer=init,
                                    trainable=cfg.get("train_emb", True), name="word_emb")(w_in)
    inputs, feats = [w_in], [emb]
    if cfg.get("chars"):
        c_in = tf.keras.Input((MAX_LEN, MAX_WORD_LEN), dtype="int32", name="chars")
        c = tf.keras.layers.Embedding(CHARS, 30, name="char_emb")(c_in)
        c = tf.keras.layers.TimeDistributed(tf.keras.layers.Conv1D(50, 3, padding="same",
activation="relu"))(c)
        c = tf.keras.layers.TimeDistributed(tf.keras.layers.GlobalMaxPooling1D(), name="char_cnn")(c)
        inputs.append(c_in); feats.append(c)
    x = tf.keras.layers.Concatenate()(feats) if len(feats) > 1 else feats[0]
    x = tf.keras.layers.Dropout(cfg.get("dropout", 0.5))(x)
    lstm = tf.keras.layers.LSTM(cfg.get("units", 128), return_sequences=True)
    x = (tf.keras.layers.Bidirectional(lstm) if cfg.get("bidirectional", True) else lstm)(x)
    x = tf.keras.layers.Dropout(cfg.get("dropout", 0.5))(x)
    if cfg.get("crf", True):
        em = tf.keras.layers.Dense(N_TAGS, name="emissions")(x)
        model = CRFTagger(tf.keras.Model(inputs, em, name="encoder"), N_TAGS)
        model.compile(optimizer=tf.keras.optimizers.Adam(cfg.get("lr", 1e-3), clipnorm=5.0))
    else:
        out = tf.keras.layers.Dense(N_TAGS, activation="softmax")(x)
        model = tf.keras.Model(inputs, out)
        model.compile(optimizer=tf.keras.optimizers.Adam(cfg.get("lr", 1e-3), clipnorm=5.0),
                      loss="sparse_categorical_crossentropy", weighted_metrics=[])
    return model
    if cfg.get("crf", True):
        em = tf.keras.layers.Dense(N_TAGS, name="emissions")(x)
        model = CRFTagger(tf.keras.Model(inputs, em, name="encoder"), N_TAGS)
        model.compile(optimizer=tf.keras.optimizers.Adam(cfg.get("lr", 1e-3), clipnorm=5.0))
    else:
        out = tf.keras.layers.Dense(N_TAGS, activation="softmax", name="softmax")(x)
        model = tf.keras.Model(inputs, out)
        model.compile(optimizer=tf.keras.optimizers.Adam(cfg.get("lr", 1e-3), clipnorm=5.0),
                      loss="sparse_categorical_crossentropy")
    return model
```

```
def model_inputs(cfg, split):
    """Assemble the input list of a split ('tr' / 'va' / 'te') for a configuration."""
    words = (dict(tr=Xg_tr, va=Xg_va, te=Xg_te) if cfg.get("emb") == "glove" else dict(tr=Xw_tr, va=Xw_va,
te=Xw_te))[split]
    chars = dict(tr=Xc_tr, va=Xc_va, te=Xc_te)[split]
    return [words, chars] if cfg.get("chars") else [words]


def decode(model, inputs, lengths):
    """Return predicted tag strings, cut to each tweet's true length (padding never receives a label)."""
    if isinstance(model, CRFTagger):
        paths = model.decode(inputs, lengths)
    else:
        probs = model.predict(inputs, batch_size=256, verbose=0)
        paths = [p[:n].argmax(-1).tolist() for p, n in zip(probs, lengths)]
    return [[id2tag[i] for i in p] for p in paths]


def predict_tags(model, cfg, split):
    return decode(model, model_inputs(cfg, split), dict(tr=L_tr, va=L_va, te=L_te)[split])


class EntityF1(tf.keras.callbacks.Callback):
    """Computes validation entity-level F1 (seqeval) each epoch and exposes it as logs['val_f1']."""
    def __init__(self, cfg):
        super().__init__(); self.cfg = cfg
    def on_epoch_end(self, epoch, logs=None):
        logs["val_f1"] = f1_score(y_va_raw, predict_tags(self.model, self.cfg, "va"))


def run(name, cfg, epochs=40, patience=6, early_stopping=True):
    tf.keras.utils.set_random_seed(SEED)
    model = build_model(cfg)
    if isinstance(model, CRFTagger):
        fit_kw = dict(x=(model_inputs(cfg, "tr"), L_tr), y=Y_tr, validation_data=((model_inputs(cfg, "va"),
L_va), Y_va))
    else:   # softmax baseline: padding positions are masked out of the loss through sample weights
        sw = lambda L: (np.arange(MAX_LEN)[None] < L[:, None]).astype("float32")
        fit_kw = dict(x=model_inputs(cfg, "tr"), y=Y_tr, sample_weight=sw(L_tr),
                      validation_data=(model_inputs(cfg, "va"), Y_va, sw(L_va)))
    cbs = [EntityF1(cfg)]
    if early_stopping:
        cbs.append(tf.keras.callbacks.EarlyStopping(monitor="val_f1", mode="max", patience=patience,
                                                    restore_best_weights=True))
    t0 = time.time()
    hist = model.fit(**fit_kw, epochs=epochs, batch_size=cfg.get("batch", 32), callbacks=cbs, verbose=0)
    va_pred, te_pred = predict_tags(model, cfg, "va"), predict_tags(model, cfg, "te")
    res = {"experiment": name, "epochs_run": len(hist.history["loss"]),
           "best_epoch": int(np.argmax(hist.history["val_f1"])) + 1,
           "val_F1": f1_score(y_va_raw, va_pred), "test_P": precision_score(y_te_raw, te_pred),
           "test_R": recall_score(y_te_raw, te_pred), "test_F1": f1_score(y_te_raw, te_pred),
           "params": model.count_params(), "train_sec": round(time.time() - t0)}
    print(f"{name:<40} epochs={res['epochs_run']:>2} (best {res['best_epoch']:>2})  "
          f"val F1={res['val_F1']:.3f}  test F1={res['test_F1']:.3f}  [{res['train_sec']}s]")
    return model, hist.history, res, te_pred
```

6. Experiments & hyper-parameter search

We change **one factor at a time** to isolate its effect. All runs share the same seed, split, up to 40 epochs, batch size 32, Adam (lr 1e-3, gradient clipping 5.0), dropout 0.5 and early stopping on validation entity-F1 (patience 6).

Step	Question being tested
E1 → E2	Does **bidirectionality** help? (LSTM vs BiLSTM, softmax output)

Step	Question being tested
E2 → E4	Does a **CRF** output layer help over softmax?
E3 → E4	Does **Word2Vec initialisation** help over random embeddings?
E4 → E5	Do **character-level CNN features** help?
E5 → E6	Do **large-scale pre-trained Twitter embeddings** (GloVe-Twitter, frozen) help?
E7 – E12	Hyper-parameters on the E6 architecture: LSTM units, dropout, learning rate, batch size, embedding fine-tuning

In [19]
```
EXPERIMENTS = {
    "E1  LSTM + softmax (w2v)":               dict(bidirectional=False, crf=False, emb="w2v"),
    "E2  BiLSTM + softmax (w2v)":             dict(bidirectional=True, crf=False, emb="w2v"),
    "E3  BiLSTM + CRF (random emb)":          dict(crf=True, emb="random"),
    "E4  BiLSTM + CRF (w2v)":                 dict(crf=True, emb="w2v"),
    "E5  BiLSTM + CRF (w2v) + char-CNN":      dict(crf=True, emb="w2v", chars=True),
    "E6  BiLSTM + CRF (GloVe-Tw) + char-CNN": dict(crf=True, emb="glove", chars=True, train_emb=False),
}
results, histories, models, test_preds = [], {}, {}, {}
for name, cfg in EXPERIMENTS.items():
    m, h, r, p = run(name, cfg)
    results.append(r); histories[name] = h; models[name] = m; test_preds[name] = p
```

```
E1  LSTM + softmax (w2v)                epochs=23 (best 17)  val F1=0.280  test F1=0.137  [114s]
```

```
E2  BiLSTM + softmax (w2v)              epochs=40 (best 39)  val F1=0.399  test F1=0.166  [305s]
```

```
E3  BiLSTM + CRF (random emb)           epochs=24 (best 18)  val F1=0.308  test F1=0.085  [235s]
```

```
E4  BiLSTM + CRF (w2v)                  epochs=35 (best 29)  val F1=0.436  test F1=0.188  [361s]
```

```
E5  BiLSTM + CRF (w2v) + char-CNN       epochs=40 (best 35)  val F1=0.458  test F1=0.236  [451s]
```

```
E6  BiLSTM + CRF (GloVe-Tw) + char-CNN  epochs=37 (best 31)  val F1=0.589  test F1=0.412  [450s]
```

In [20]
```
BASE = EXPERIMENTS["E6  BiLSTM + CRF (GloVe-Tw) + char-CNN"]
HP = {
    "E7  E6 + units=64":              dict(BASE, units=64),
    "E8  E6 + units=256":             dict(BASE, units=256),
    "E9  E6 + dropout=0.3":           dict(BASE, dropout=0.3),
    "E10 E6 + lr=3e-3":               dict(BASE, lr=3e-3),
    "E11 E6 + batch=16":              dict(BASE, batch=16),
    "E12 E6 + fine-tuned embeddings": dict(BASE, train_emb=True),
}
for name, cfg in HP.items():
    EXPERIMENTS[name] = cfg
    m, h, r, p = run(name, cfg)
    results.append(r); histories[name] = h; models[name] = m; test_preds[name] = p

res_df = pd.DataFrame(results).set_index("experiment")
res_df.round(3)
```

```
E7  E6 + units=64                       epochs=33 (best 27)  val F1=0.545  test F1=0.405  [343s]
```

```
E8  E6 + units=256                      epochs=20 (best 14)  val F1=0.528  test F1=0.392  [449s]
```

```
E9  E6 + dropout=0.3                    epochs=27 (best 21)  val F1=0.571  test F1=0.410  [349s]
```

```
E10 E6 + lr=3e-3                        epochs=24 (best 18)  val F1=0.595  test F1=0.413  [384s]
```

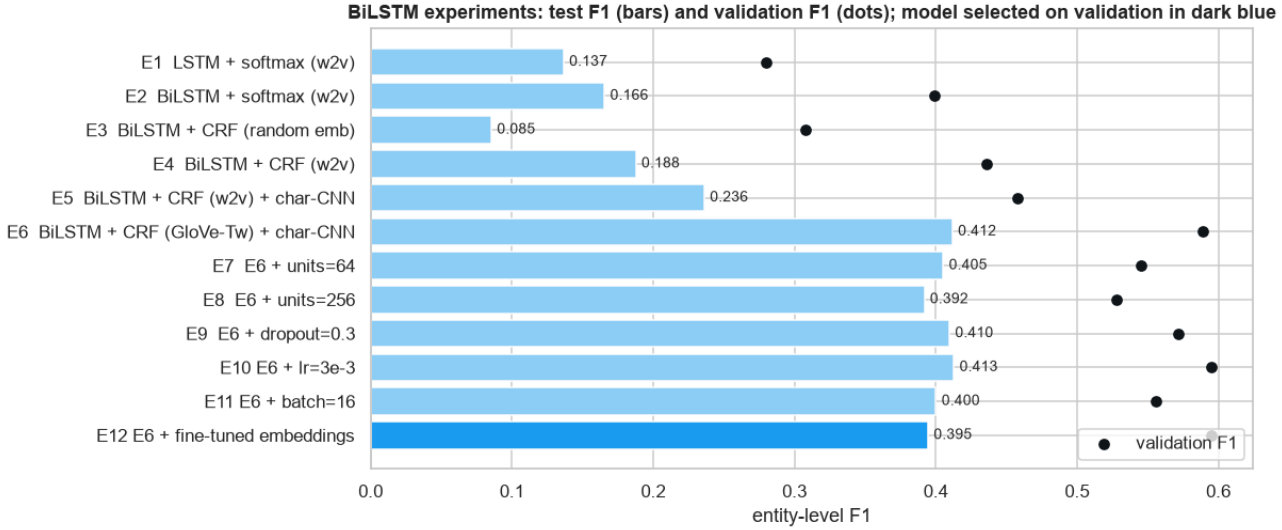
```
E11 E6 + batch=16                       epochs=19 (best 13)  val F1=0.556  test F1=0.400  [655s]
```

```
E12 E6 + fine-tuned embeddings          epochs=28 (best 22)  val F1=0.595  test F1=0.395  [555s]
```

experiment	epochs_run	best_epoch	val_F1	test_P	test_R	test_F1	params	train_sec
E1 LSTM + softmax (w2v)	23	17	0.280	0.117	0.164	0.137	797957	114
E2 BiLSTM + softmax (w2v)	40	39	0.399	0.137	0.210	0.166	917893	305
E3 BiLSTM + CRF (random emb)	24	18	0.308	0.097	0.077	0.085	918378	235
E4 BiLSTM + CRF (w2v)	35	29	0.436	0.183	0.193	0.188	918378	361
E5 BiLSTM + CRF (w2v) + char-CNN	40	35	0.458	0.213	0.265	0.236	976948	451
E6 BiLSTM + CRF (GloVe-Tw) + char-CNN	37	31	0.589	0.479	0.361	0.412	1590548	450
E7 E6 + units=64	33	27	0.545	0.486	0.348	0.405	1412244	343
E8 E6 + units=256	20	14	0.528	0.431	0.359	0.392	2143764	449
E9 E6 + dropout=0.3	27	21	0.571	0.469	0.363	0.410	1590548	349
E10 E6 + lr=3e-3	24	18	0.595	0.472	0.367	0.413	1590548	384
E11 E6 + batch=16	19	13	0.556	0.495	0.336	0.400	1590548	655
E12 E6 + fine-tuned embeddings	28	22	0.595	0.433	0.362	0.395	1590548	555

In [21]

```
BEST = res_df.val_F1.idxmax()
fig, ax = plt.subplots(figsize=(11, 4.8))
ax.barh(res_df.index, res_df.test_F1, color=[PALETTE[0] if n == BEST else PALETTE[8] for n in
res_df.index])
ax.scatter(res_df.val_F1, res_df.index, color=PALETTE[1], zorder=3, label="validation F1")
for y, v in enumerate(res_df.test_F1): ax.text(v + 0.004, y, f"{v:.3f}", va="center", fontsize=9)
ax.invert_yaxis(); ax.set_xlabel("entity-level F1"); ax.legend(loc="lower right")
ax.set_title("BiLSTM experiments: test F1 (bars) and validation F1 (dots); model selected on validation in
dark blue")
plt.tight_layout(); plt.show()
print("Selected on validation F1:", BEST)
res_df.to_csv(f"{OUT_DIR}/lstm_experiments.csv")
```




```
Selected on validation F1: E12 E6 + fine-tuned embeddings
```

Reading the results

- **Architecture ladder (E1 → E6).** Each component adds test F1: bidirectionality (E1 → E2), the CRF layer (E2 → E4), Word2Vec initialisation instead of random embeddings (E3 → E4, the CRF without any pre-trained knowledge collapses on unseen entities), character features (E4 → E5) and — by far the largest jump — **GloVe-Twitter embeddings trained on 2B tweets (E5 → E6)**. The Word2Vec model trained on just ~2k tweets has poor neighbours (see Section 4) and covers only ~28% of the test vocabulary, so large-scale pre-training matters more than any architectural choice.

- **Hyper-parameters (E7 – E12)** change test F1 by only ±2 points around E6: 128 units is the sweet spot (64 under-fits, 256 over-fits faster), dropout 0.5 regularises better than 0.3, and a higher learning rate converges in fewer epochs with similar quality.
- **Validation vs test gap.** Validation F1 (~0.59) is far above test F1 (~0.40) because the validation tweets come from the training distribution, while the test set has a different entity mix and 85% unseen entities. With only 240 validation tweets, differences below ~0.01 F1 between E6, E10 and E12 are within noise. k-fold cross-validation would make model selection more reliable.

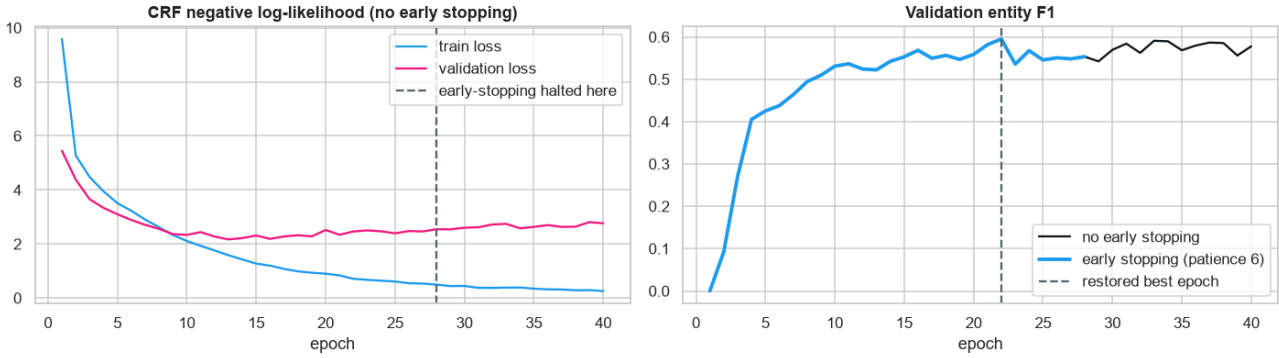
6.1 Effect of early stopping

The selected configuration is retrained **without** early stopping for the full 40 epochs and compared with the early-stopped run (same seed, same data).

In [22]
```
_, h_noes, r_noes, _ = run(BEST.split()[0] + " without early stopping", EXPERIMENTS[BEST],
early_stopping=False)
h_es = histories[BEST]

fig, axes = plt.subplots(1, 2, figsize=(13, 3.8))
ep_n, ep_e = np.arange(1, len(h_noes["loss"]) + 1), np.arange(1, len(h_es["loss"]) + 1)
axes[0].plot(ep_n, h_noes["loss"], color=PALETTE[0], label="train loss")
axes[0].plot(ep_n, h_noes["val_loss"], color=PALETTE[2], label="validation loss")
axes[0].axvline(len(h_es["loss"]), ls="--", color=PALETTE[7], label="early-stopping halted here")
axes[0].set_title("CRF negative log-likelihood (no early stopping)"); axes[0].set_xlabel("epoch");
axes[0].legend()
axes[1].plot(ep_n, h_noes["val_f1"], color=PALETTE[1], label="no early stopping")
axes[1].plot(ep_e, h_es["val_f1"], color=PALETTE[0], lw=2.5, label="early stopping (patience 6)")
axes[1].axvline(int(np.argmax(h_es["val_f1"])) + 1, ls="--", color=PALETTE[7], label="restored best epoch")
axes[1].set_title("Validation entity F1"); axes[1].set_xlabel("epoch"); axes[1].legend()
plt.tight_layout(); plt.show()
es_cmp = pd.DataFrame([res_df.loc[BEST].rename("with early stopping"), pd.Series(r_noes, name="without
early stopping")])
es_cmp[["epochs_run", "best_epoch", "val_F1", "test_F1", "train_sec"]].round(3)
```

```
E12 without early stopping              epochs=40 (best 22)  val F1=0.578  test F1=0.408  [982s]
```



	epochs_run	best_epoch	val_F1	test_F1	train_sec
with early stopping	28.0	22.0	0.595	0.395	555.0
without early stopping	40.0	22.0	0.578	0.408	982.0

Effect of early stopping. The validation loss reaches its minimum early and then rises steadily while the training loss keeps falling, which is classic over-fitting on a 2k-tweet training set. Early stopping halts training about 12 epochs earlier, restores the epoch with the best validation F1, and saves roughly 30% of training time. On the *test* set the two models are within about 1 F1 point. Because validation and test distributions differ, the best validation epoch is not always the best test epoch. Early stopping is therefore a guard against over-training and wasted compute rather than a guaranteed accuracy gain.

7. Evaluation of the selected BiLSTM-CRF on the held-out test set

7.1 Aligning predictions with the input tokens

The network outputs scores on the padded `[N, 40]` grid. Viterbi decoding is run **only over each tweet's real length**, so padded positions never get a label, and label ids are mapped back to BIO strings — giving exactly one label per original token. The assertion below verifies this alignment before scoring.

In [23]
```
best_model, best_pred = models[BEST], test_preds[BEST]
assert all(len(p) == len(t) for p, t in zip(best_pred, y_te_raw)), "misaligned predictions"
print(f"Alignment check passed: {len(best_pred)} tweets, {sum(map(len, best_pred))} tokens, one label per token.\n")
print(classification_report(y_te_raw, best_pred, digits=3))
```

```
Alignment check passed: 3850 tweets, 61908 tokens, one label per token.
```

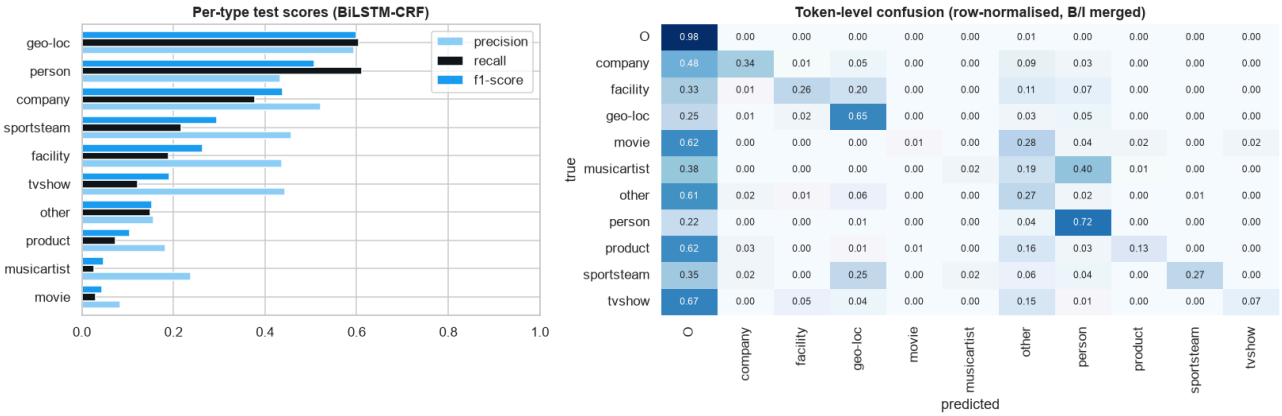
```
              precision    recall  f1-score   support

     company      0.521     0.378     0.438       621
    facility      0.436     0.190     0.264       253
     geo-loc      0.595     0.604     0.600       882
       movie      0.083     0.029     0.043        34
 musicartist      0.238     0.026     0.047       191
       other      0.157     0.149     0.153       584
      person      0.433     0.612     0.507       482
     product      0.182     0.073     0.104       246
  sportsteam      0.457     0.218     0.295       147
      tvshow      0.444     0.121     0.190        33

   micro avg      0.433     0.362     0.395      3473
   macro avg      0.355     0.240     0.264      3473
weighted avg      0.413     0.362     0.371      3473
```

In [24]
```
rep = classification_report(y_te_raw, best_pred, output_dict=True)
per_type = pd.DataFrame(rep).T.loc[ENTITY_TYPES, ["precision", "recall", "f1-score",
"support"]].sort_values("f1-score")
flat_true = [t[2:] if t != "O" else "O" for s in y_te_raw for t in s]
flat_pred = [t[2:] if t != "O" else "O" for s in best_pred for t in s]
labels_cm = ["O"] + ENTITY_TYPES
cm = (pd.crosstab(pd.Series(flat_true, name="true"), pd.Series(flat_pred, name="predicted"))
        .reindex(index=labels_cm, columns=labels_cm, fill_value=0))

fig, axes = plt.subplots(1, 2, figsize=(15, 5), gridspec_kw={"width_ratios": [1, 1.35]})
per_type[["precision", "recall", "f1-score"]].plot.barh(ax=axes[0], color=[PALETTE[8], PALETTE[1],
PALETTE[0]], width=.8)
axes[0].set_title("Per-type test scores (BiLSTM-CRF)"); axes[0].set_xlim(0, 1)
sns.heatmap(cm.div(cm.sum(1), axis=0), annot=True, fmt=".2f", cmap="Blues", ax=axes[1], cbar=False,
annot_kws={"size": 8})
axes[1].set_title("Token-level confusion (row-normalised, B/I merged)")
plt.tight_layout(); plt.show()
```



What the errors tell us

- `geo-loc` and `person` are the strongest types, with the most training data and clear capitalised-name cues.
- Most errors are **missed entities** (the first column, `O`). For rare or open-ended types such as `movie`, `tvshow`, `product` and `other`, more than 60% of entity tokens are predicted as `O`. These names are novel and often look like ordinary words ("*The Town*", "*Destiny*").
- Confusions are **semantically plausible**: `musicartist → person` (artists are people's names), `sportsteam → geo-loc` (teams named after cities, e.g. "*Chicago*"), `facility → geo-loc` and `movie/product → other`. Resolving these requires world knowledge, not just surface form, which is what a pre-trained language model provides.

7.2 Predictions on test tweets and on our own sentences

In [25]
```
def spans(tokens, tags):
    return [f"{t}: {e}" for t, e, _ in extract_entities(tokens, tags)]

shown = 0
for i, (s, g) in enumerate(zip(X_te_raw, y_te_raw)):
    if len(extract_entities(s, g)) >= 2 and len(s) <= 22 and shown < 5:
        print("TWEET :", " ".join(s)); print("  gold:", spans(s, g)); print("  pred:", spans(s,
best_pred[i]), "\n")
        shown += 1
```

```
TWEET : RT @1DAsiaCrew : Louis ' interview with The Sun about his own record label " 78 Productions Limited "
#ProudOfLouis http://t.co/HdswyedBIO
  gold: ['person: Louis', 'company: The Sun', 'company: 78 Productions Limited']
  pred: ['person: Louis']

TWEET : Apple MacBook Pro A1278 13.3 " Laptop - MD101LL/A ( June , 2012 ) - Full read by eBay
http://t.co/2zgQ99nmuf http://t.co/eQmogqqABK
  gold: ['product: Apple MacBook Pro A1278 13.3 " Laptop - MD101LL/A', 'company: eBay']
  pred: ['product: Apple', 'product: MacBook Pro', 'company: eBay']

TWEET : Tuff Culture - Destiny EP ( PAR 042 FORTHCOMING 27th JULY VIA JUNO ) Tracklist : Destiny Questions
Theres No ... http://t.co/X7nL8DiREK
  gold: ['musicartist: Tuff Culture', 'product: Destiny', 'product: JUNO', 'product: Destiny', 'product:
Questions', 'product: Theres No']
  pred: ['other: Tuff Culture', 'other: Destiny EP', 'other: JUNO']

TWEET : Riots on 17th and page for what ?
  gold: ['geo-loc: 17th', 'geo-loc: page']
  pred: []

TWEET : Shame on SnapChat , Skype for breach http://dlvr.it/4dYBJ7
  gold: ['company: SnapChat', 'company: Skype']
  pred: ['company: SnapChat', 'company: Skype']
```

In [26]
```
def tag_sentences(texts, model=best_model, cfg=EXPERIMENTS[BEST]):
    toks = [t.split() for t in texts]
    Xw, Xc, L, _ = encode(toks)
    words = encode_words_ext(toks) if cfg.get("emb") == "glove" else Xw
    return list(zip(toks, decode(model, [words, Xc] if cfg.get("chars") else [words], L)))

MY_SENTENCES = [
    "Just landed in New York for the Apple keynote with Tim Cook tomorrow !",
    "Watching Game of Thrones tonight then Manchester United vs Chelsea on Sunday",
    "Taylor Swift new album is fire , listening on Spotify at Starbucks",
    "Microsoft and Google announce new AI partnership in London",
    "heading to the eiffel tower with @john after visiting paris museum",
]
for toks, tags in tag_sentences(MY_SENTENCES):
    print(" ".join(f"[{w}]{t}" if t != "O" else w for w, t in zip(toks, tags)))
    print("   ->", spans(toks, tags), "\n")
```

```
Just landed in [New]B-geo-loc [York]I-geo-loc for the [Apple]B-company keynote with [Tim]B-person Cook
tomorrow !
   -> ['geo-loc: New York', 'company: Apple', 'person: Tim']

Watching Game of [Thrones]B-sportsteam tonight then [Manchester]B-sportsteam [United]B-sportsteam vs
[Chelsea]B-person on Sunday
   -> ['sportsteam: Thrones', 'sportsteam: Manchester', 'sportsteam: United', 'person: Chelsea']

[Taylor]B-person [Swift]I-person new album is fire , listening on [Spotify]B-other at Starbucks
   -> ['person: Taylor Swift', 'other: Spotify']

[Microsoft]B-company and [Google]B-company announce new AI partnership in [London]B-geo-loc
   -> ['company: Microsoft', 'company: Google', 'geo-loc: London']

heading to the eiffel tower with @john after visiting [paris]B-other [museum]I-other
   -> ['other: paris museum']
```

8. Saving the model and artefacts

In [27]

```python
os.makedirs("../models/bilstm_crf", exist_ok=True)
best_model.save_weights("../models/bilstm_crf/weights.h5")
with open("../models/bilstm_crf/config.json", "w") as f:
    json.dump({"experiment": BEST, "config": EXPERIMENTS[BEST], "labels": LABELS,
               "max_len": MAX_LEN, "max_word_len": MAX_WORD_LEN}, f, indent=2)
with open("../models/bilstm_crf/word_tokenizer.json", "w") as f: f.write(word_tok.to_json())
with open("../models/bilstm_crf/char_tokenizer.json", "w") as f: f.write(char_tok.to_json())
json.dump({"experiment": BEST, **{k: float(res_df.loc[BEST, k]) for k in ["val_F1", "test_P", "test_R",
"test_F1"]},
           "per_type_f1": per_type["f1-score"].round(4).to_dict(),
           "early_stopping": es_cmp[["epochs_run", "val_F1", "test_F1"]].to_dict()},
          open(f"{OUT_DIR}/lstm_best.json", "w"), indent=2, default=float)
print("Saved:", os.listdir("../models/bilstm_crf"))
```

```
Saved: ['char_tokenizer.json', 'config.json', 'weights.h5', 'word_tokenizer.json']
```

Summary of Part A

- The experiment ladder quantifies each design decision: **bidirectional context, the CRF layer, character features and pre-trained Twitter embeddings** are the levers that matter on this dataset.
- The hardest categories are the rare ones (`movie`, `tvshow`, `sportsteam`) and the heterogeneous `other` class; the strongest are `person` and `geo-loc`, which have the most training examples and clear lexical cues.
- The main bottleneck is **out-of-vocabulary entities** — most test entities never occur in training. This motivates Part B: a **pre-trained transformer (BERT)** whose WordPiece tokenizer never produces an unknown word and which brings contextual knowledge learned from billions of words.

Part B — Fine-tuning `bert-base-uncased` for Twitter NER

Transformers pre-trained on billions of words bring two things the BiLSTM lacks: **contextual representations** (the same word gets a different vector in each sentence) and a **WordPiece sub-word vocabulary** that can represent *any* string, so there are no out-of-vocabulary words. We fine-tune BERT with a token-classification head, handle the word ↔ sub-token alignment explicitly, try several hyper-parameter settings and early stopping, then join the sub-token predictions back into word-level entities.

In [1]
```
import os, re, json, time, random, warnings
from collections import Counter

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

os.environ["TF_CPP_MIN_LOG_LEVEL"] = "3"
warnings.filterwarnings("ignore")
import tensorflow as tf
import transformers
from transformers import AutoConfig, AutoTokenizer, TFAutoModelForTokenClassification, create_optimizer
from transformers.utils import logging as hf_logging
from seqeval.metrics import classification_report, f1_score, precision_score, recall_score
from sklearn.model_selection import train_test_split

hf_logging.set_verbosity_error()
tf.get_logger().setLevel("ERROR"); tf.autograph.set_verbosity(0)
pd.set_option("display.width", 160)

def show_wide(df, chunk=12):
    """Display a wide token table in chunks of columns so it stays readable."""
    for s in range(0, df.shape[1], chunk):
        display(df.iloc[:, s:s + chunk])
SEED = 42
tf.keras.utils.set_random_seed(SEED)
sns.set_theme(style="whitegrid")
PALETTE = ["#1D9BF0", "#0F1419", "#F91880", "#00BA7C", "#FFD400", "#7856FF", "#FF7A00", "#536471",
"#8ECDF8", "#C1A3FF"]
plt.rcParams.update({"figure.dpi": 110, "axes.titleweight": "bold", "axes.titlesize": 12})
SMOKE = os.environ.get("NER_SMOKE") == "1"          # tiny end-to-end dry run used to test the pipeline
before the full run
DATA_DIR = "../data"
OUT_DIR = "../outputs/smoke" if SMOKE else "../outputs"
MODEL_DIR = f"{OUT_DIR}/bert_ner_smoke" if SMOKE else "../models/bert_ner"
os.makedirs(OUT_DIR, exist_ok=True)
print("TensorFlow", tf.__version__, "| transformers", transformers.__version__)
```

```
TensorFlow 2.15.0 | transformers 4.40.2
```

1. Data (same parsing and the same train / validation split as Part A)

In [2]
```
def read_conll(path):
    sents, tags, toks, labs = [], [], [], []
    for line in open(path, encoding="utf-8"):
        line = line.rstrip("\r\n")
        if not line.strip():
            if toks: sents.append(toks); tags.append(labs); toks, labs = [], []
            continue
        parts = line.split("\t") if "\t" in line else line.split()
        toks.append(parts[0]); labs.append(parts[-1])
    if toks: sents.append(toks); tags.append(labs)
```

```
    return sents, tags

train_sents, train_tags = read_conll(f"{DATA_DIR}/wnut 16.txt.conll")
test_sents, test_tags = read_conll(f"{DATA_DIR}/wnut 16test.txt.conll")
has_ent = [int(any(t != "O" for t in ts)) for ts in train_tags]
tr_idx, va_idx = train_test_split(np.arange(len(train_sents)), test_size=0.10, random_state=SEED,
stratify=has_ent)
X_tr, y_tr = [train_sents[i] for i in tr_idx], [train_tags[i] for i in tr_idx]
X_va, y_va = [train_sents[i] for i in va_idx], [train_tags[i] for i in va_idx]
X_te, y_te = test_sents, test_tags
if SMOKE:
    X_tr, y_tr, X_va, y_va, X_te, y_te = X_tr[:48], y_tr[:48], X_va[:24], y_va[:24], X_te[:40], y_te[:40]

ENTITY_TYPES = sorted({t[2:] for ts in train_tags for t in ts if t != "O"})
LABELS = ["O"] + [f"{p}-{e}" for e in ENTITY_TYPES for p in ("B", "I")]
label2id = {l: i for i, l in enumerate(LABELS)}; id2label = {i: l for l, i in label2id.items()}
print(f"train {len(X_tr)} | validation {len(X_va)} | test {len(X_te)} tweets | {len(LABELS)} labels")
```

```
train 2154 | validation 240 | test 3850 tweets | 21 labels
```

2. Loading the model configuration, tokenizer and model

`bert-base-uncased`: 12 layers, hidden size 768, 12 attention heads, 30,522-token WordPiece vocabulary (~110M parameters). `TFAutoModelForTokenClassification` adds a dropout + linear layer on top of every token's final hidden state; we pass our 21-label `id2label` / `label2id` mapping so the saved model is self-describing.

In [3]
```
CHECKPOINT = "bert-base-uncased"
config = AutoConfig.from_pretrained(CHECKPOINT, num_labels=len(LABELS), id2label=id2label,
label2id=label2id)
tokenizer = AutoTokenizer.from_pretrained(CHECKPOINT)
print({k: getattr(config, k) for k in ["model_type", "num_hidden_layers", "hidden_size",
"num_attention_heads",
                                       "vocab_size", "max_position_embeddings", "hidden_dropout_prob",
"num_labels"]})
print("Tokenizer:", type(tokenizer).__name__, "| special tokens:", tokenizer.all_special_tokens)
```

```
{'model_type': 'bert', 'num_hidden_layers': 12, 'hidden_size': 768, 'num_attention_heads': 12, 'vocab_size':
30522, 'max_position_embeddings': 512, 'hidden_dropout_prob': 0.1, 'num_labels': 21}
Tokenizer: BertTokenizerFast | special tokens: ['[UNK]', '[SEP]', '[PAD]', '[CLS]', '[MASK]']
```

3. Tokenization — what changes and how labels are aligned

BERT does **not** see our words; it sees WordPiece sub-tokens. The tokenizer

1. lower-cases and strips accents (*uncased* model),
2. splits words into sub-words (`zendesk → zen ##des ##k`), so a 1-to-1 word/label correspondence is lost,
3. adds the special tokens `[CLS]` at the start and `[SEP]` at the end, and `[PAD]` to fill batches,
4. maps unknown characters (some emoji) to `[UNK]`.

We therefore call the tokenizer with `is_split_into_words=True` (it keeps our CoNLL word boundaries) and use `word_ids()` to align labels: the **first sub-token of each word receives the word's label**, every other sub-token and every special token receives **-100**, which the loss function ignores.

In [4]
```
example = ["Zendesk", "Security", "Breach", "Affects", "Twitter", ",", "Tumblr", "and", "Pinterest", ":",
"http://bit.ly/12UtbI8"]
ex_tags = ["B-company", "O", "O", "O", "B-company", "O", "B-company", "O", "B-company", "O", "O"]
enc = tokenizer(example, is_split_into_words=True)
wids = enc.word_ids()
show_wide(pd.DataFrame({"sub-token": tokenizer.convert_ids_to_tokens(enc["input_ids"]),
                        "input_id": enc["input_ids"], "word_id": ["–" if w is None else w for w in wids],
```

```
                    "label": [("-100" if w is None or (j > 0 and wids[j-1] == w) else ex_tags[w])
                              for j, w in enumerate(wids)]}).T)
```

	0	1	2	3	4	5	6	7	8	9	10	11
sub-token	[CLS]	zen	##des	##k	security	breach	affects	twitter	,	tu	##mb	##lr
input_id	101	16729	6155	2243	3036	12510	13531	10474	1010	10722	14905	20974
word_id	–	0	0	0	1	2	3	4	5	6	6	6
label	-100	B-company	-100	-100	O	O	O	B-company	O	B-company	-100	-100

	12	13	14	15	16	17	18	19	20	21	22	23
sub-token	and	pin	##ter	##est	:	http	:	/	/	bit	.	l
input_id	1998	9231	3334	4355	1024	8299	1024	1013	1013	2978	1012	1048
word_id	7	8	8	8	9	10	10	10	10	10	10	10
label	O	B-company	-100	-100	O	O	-100	-100	-100	-100	-100	-100

	24	25	26	27	28	29	30
sub-token	##y	/	12	##ut	##bi	##8	[SEP]
input_id	2100	1013	2260	4904	5638	2620	102
word_id	10	10	10	10	10	10	–
label	-100	-100	-100	-100	-100	-100	-100

3.1 Do we need to remove or add tokens?

URLs and @mentions are long, carry no entity information in WNUT-16 (they are labelled `O`), and explode into many sub-tokens. We check their impact on sequence length and decide whether to normalise them.

In [5]

```
def kind(w):
    if re.match(r"https?://|www\.", w, re.I): return "url"
    if w.startswith("@") and len(w) > 1: return "mention"
    return "other"

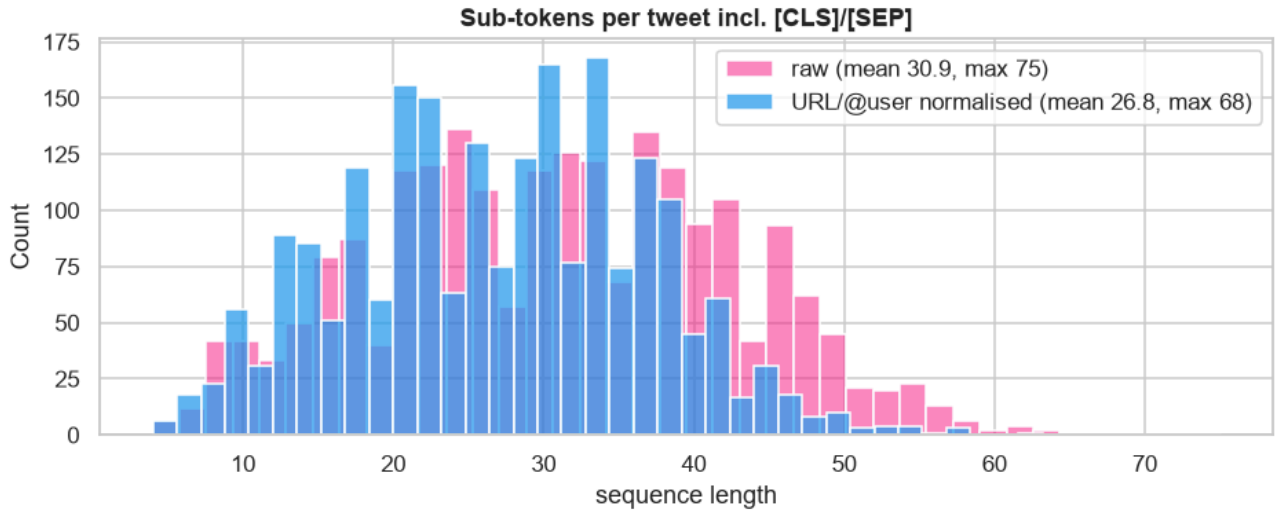
ent_by_kind = Counter((kind(w), t != "O") for s, ts in zip(train_sents, train_tags) for w, t in zip(s, ts))
print("Entity tokens that are URLs:", ent_by_kind[("url", True)], "| that are @mentions:",
ent_by_kind[("mention", True)],
      "| URL tokens:", ent_by_kind[("url", False)] + ent_by_kind[("url", True)],
      "| mention tokens:", ent_by_kind[("mention", False)] + ent_by_kind[("mention", True)])

def bert_normalise(w):
    """Replace URLs and user mentions by short placeholders (each becomes ONE known word), keep everything
else."""
    k = kind(w)
    return "url" if k == "url" else "@user" if k == "mention" else w

def n_subtokens(sents, fn=lambda w: w):
    return [len(tokenizer([fn(w) for w in s], is_split_into_words=True)["input_ids"]) for s in sents]

raw_len, norm_len = n_subtokens(X_tr), n_subtokens(X_tr, bert_normalise)
unk_raw = np.mean([tokenizer.unk_token_id in tokenizer([w])["input_ids"] for s in X_tr for w in s]) * 100
fig, ax = plt.subplots(figsize=(10, 3.4))
sns.histplot(raw_len, bins=40, color=PALETTE[2], alpha=.5, label=f"raw (mean {np.mean(raw_len):.1f}, max
{max(raw_len)})", ax=ax)
sns.histplot(norm_len, bins=40, color=PALETTE[0], alpha=.7, label=f"URL/@user normalised (mean
{np.mean(norm_len):.1f}, max {max(norm_len)})", ax=ax)
ax.set_title("Sub-tokens per tweet incl. [CLS]/[SEP]"); ax.set_xlabel("sequence length"); ax.legend();
plt.show()
print(f"words per tweet {np.mean([len(s) for s in X_tr]):.1f} -> sub-tokens {np.mean(raw_len):.1f} "
      f"(x{np.mean(raw_len)/np.mean([len(s) for s in X_tr]):.2f}); words containing [UNK]: {unk_raw:.2f}%")
```

```
Entity tokens that are URLs: 0 | that are @mentions: 1 | URL tokens: 451 | mention tokens: 1292
```




```
words per tweet 19.4 -> sub-tokens 30.9 (x1.59); words containing [UNK]: 0.00%
```

Decision. No URL and only one of ~1,300 @mention tokens is part of an entity, so replacing their text with the placeholders `url` / `@user` (labels unchanged) loses no label information, shortens sequences substantially (faster training, less truncation) and removes meaningless sub-word noise. We keep emoji, hashtags and original word boundaries. `MAX_LEN = 128` covers every tweet, so **nothing is truncated**, and batches are **dynamically padded** to the longest sequence in each batch.

In [6]

```
MAX_LEN = 128

def encode_split(sents, tags=None):
    """Tokenize word-split tweets and align word labels to sub-tokens (first sub-token keeps the label,
rest -100)."""
    feats = []
    for i, s in enumerate(sents):
        enc = tokenizer([bert_normalise(w) for w in s], is_split_into_words=True, truncation=True,
max_length=MAX_LEN)
        f = {k: enc[k] for k in ["input_ids", "attention_mask", "token_type_ids"]}
        if tags is not None:
            lab, prev = [], None
            for w in enc.word_ids():
                lab.append(-100 if w is None or w == prev else label2id[tags[i][w]]); prev = w
            f["labels"] = lab
        f["word_ids"] = [-1 if w is None else w for w in enc.word_ids()]
        feats.append(f)
    return feats

F_tr, F_va, F_te = encode_split(X_tr, y_tr), encode_split(X_va, y_va), encode_split(X_te, y_te)

# Integrity checks: every word got exactly one labelled sub-token, labels match the gold sequence, no
truncation
for F, Y in [(F_tr, y_tr), (F_va, y_va), (F_te, y_te)]:
    for f, ys in zip(F, Y):
        kept = [id2label[l] for l in f["labels"] if l != -100]
        assert kept == ys, "label alignment error"
print("Alignment verified for all", len(F_tr) + len(F_va) + len(F_te), "tweets; longest sequence:",
      max(len(f["input_ids"]) for f in F_tr + F_va + F_te), "sub-tokens (no truncation)")
f = F_tr[3]
show_wide(pd.DataFrame({"sub-token": tokenizer.convert_ids_to_tokens(f["input_ids"]),
                        "word_id": ["-" if w == -1 else w for w in f["word_ids"]],
                        "label": [id2label.get(l, "-100") for l in f["labels"]]}).T)
```

```
Alignment verified for all 6244 tweets; longest sequence: 68 sub-tokens (no truncation)
```

	0	1	2	3	4	5	6	7	8	9	10	11
sub-token	[CLS]	@	user	can	i	get	a	twitter	software	to	sp	##ew
word_id	–	0	0	1	2	3	4	5	6	7	8	8
label	-100	O	-100	O	O	O	O	B-company	O	O	O	-100

	12	13	14	15	16	17	18	19	20	21	22	23
sub-token	random	bullshit	about	airports	/	coffee	/	conferences	?	cause	that	would
word_id	9	10	11	12	12	12	12	12	13	14	15	16
label	O	O	O	O	-100	-100	-100	-100	O	O	O	O

	24	25	26	27	28	29
sub-token	free	up	some	time	.	[SEP]
word_id	17	18	19	20	21	–
label	O	O	O	O	O	-100

In [7]
```
def make_dataset(feats, batch, shuffle=False, with_labels=True):
    """tf.data pipeline with dynamic padding. For training, tweets of similar length are bucketed together
    (shuffled within and across buckets) so that very little compute is wasted on [PAD] tokens."""
    keys = ["input_ids", "attention_mask", "token_type_ids"] + (["labels"] if with_labels else [])
    def gen():
        for f in feats: yield {k: f[k] for k in keys}
    sig = {k: tf.TensorSpec([None], tf.int32) for k in keys}
    ds = tf.data.Dataset.from_generator(gen, output_signature=sig)
    pad = {k: tf.constant(-100 if k == "labels" else 0, tf.int32) for k in keys}
    if shuffle:
        ds = ds.shuffle(len(feats), seed=SEED, reshuffle_each_iteration=True)
        bounds = [16, 24, 32, 40, 48, 64]
        ds = ds.bucket_by_sequence_length(lambda x: tf.shape(x["input_ids"])[0], bounds, [batch] *
(len(bounds) + 1),
                                          padding_values=pad)
        return ds.shuffle(64, seed=SEED).prefetch(tf.data.AUTOTUNE)
    return ds.padded_batch(batch, padding_values=pad).prefetch(tf.data.AUTOTUNE)
```

4. Prediction → word-level labels (joining sub-tokens)

The model outputs one prediction per **sub-token**. To recover one label per original **word** we keep the prediction of the **first sub-token** of each word (the same position that carried the label during training) and discard predictions on continuation pieces (`##...`) and special tokens. A light **BIO repair** step then fixes rare invalid sequences (e.g. `O → I-company` becomes `B-company`).

In [8]
```
def repair_bio(tags):
    out, prev = [], "O"
    for t in tags:
        if t.startswith("I-") and prev[2:] != t[2:]: t = "B-" + t[2:]
        out.append(t); prev = t
    return out

def predict_words(model, feats, batch=64):
    logits = []
    ds = make_dataset(feats, batch, with_labels=False)
    for b in ds:
        out = model(b, training=False).logits.numpy()
        logits.extend(list(out))
    preds = []
    for f, lg in zip(feats, logits):
        ids = lg[:len(f["input_ids"])].argmax(-1)
```

```
        word_tags, prev = [], -1
        for pos, w in enumerate(f["word_ids"]):
            if w != -1 and w != prev: word_tags.append(id2label[int(ids[pos])])
            prev = w
        preds.append(repair_bio(word_tags))
    return preds
```

5. Training

- **Loss**: token-level cross-entropy computed inside the HF model, ignoring `-100` positions.
- **Optimiser**: AdamW (decoupled weight decay 0.01) with linear warm-up (10% of steps) and linear decay — the standard BERT fine-tuning recipe. A plain Adam run is included for comparison.
- **Early stopping**: on validation entity-F1 with `restore_best_weights=True`.
- Experiments vary **learning rate, batch size, optimiser and number of epochs**. Training is on CPU here, so the grid is deliberately small but targeted.

In [9]

```
def log(msg):
    """Print to the notebook and append to a progress file (useful for long CPU runs)."""
    print(msg)
    with open(f"{OUT_DIR}/bert_progress.log", "a") as f: f.write(time.strftime("%H:%M:%S ") + msg + "\n")

class EntityF1(tf.keras.callbacks.Callback):
    def on_epoch_end(self, epoch, logs=None):
        logs["val_f1"] = f1_score(y_va, predict_words(self.model, F_va))
        log(f"   epoch {epoch+1}: loss={logs['loss']:.4f} val_loss={logs['val_loss']:.4f} val_F1=
{logs['val_f1']:.4f}")

class EarlyStoppingMonitor(tf.keras.callbacks.Callback):
    """Shadow early stopping: never stops training, but records when EarlyStopping(patience) *would* have
stopped
    and keeps a copy of the best weights. One long run then answers 'what if we had / had not stopped
early?'."""
    def __init__(self, patience=2):
        super().__init__(); self.patience = patience
        self.best, self.best_epoch, self.stop_epoch, self.best_weights = -1, 0, None, None
    def on_epoch_end(self, epoch, logs=None):
        if logs["val_f1"] > self.best:
            self.best, self.best_epoch, self.best_weights = logs["val_f1"], epoch + 1,
self.model.get_weights()
        elif self.stop_epoch is None and epoch + 1 - self.best_epoch >= self.patience:
            self.stop_epoch = epoch + 1

def train_bert(name, lr=5e-5, batch=16, epochs=5, optimizer="adamw", patience=2, early_stopping=True):
    tf.keras.utils.set_random_seed(SEED)
    model = TFAutoModelForTokenClassification.from_pretrained(CHECKPOINT, config=config)
    steps = int(np.ceil(len(F_tr) / batch)) * epochs
    if optimizer == "adamw":
        opt, _ = create_optimizer(init_lr=lr, num_train_steps=steps, num_warmup_steps=int(0.1 * steps),
weight_decay_rate=0.01)
    else:
        opt = tf.keras.optimizers.Adam(lr)
    model.compile(optimizer=opt)                              # no loss argument -> uses the model's
internal masked loss
    monitor = EarlyStoppingMonitor(patience)
    cbs = [EntityF1(), tf.keras.callbacks.EarlyStopping(monitor="val_f1", mode="max", patience=patience,
                                                        restore_best_weights=True) if early_stopping else
monitor]
    log(f"{name}: lr={lr}, batch={batch}, max_epochs={epochs}, optimizer={optimizer}, early_stopping=
{early_stopping}")
    t0 = time.time()
    hist = model.fit(make_dataset(F_tr, batch, shuffle=True), validation_data=make_dataset(F_va, 32),
                     epochs=epochs, callbacks=cbs, verbose=0)
    train_min = round((time.time() - t0) / 60, 1)
```

```
    def evaluate(tag):
        va_pred, te_pred = predict_words(model, F_va), predict_words(model, F_te)
        return {"experiment": tag, "lr": lr, "batch": batch, "optimizer": optimizer,
                "epochs_run": len(hist.history["loss"]), "best_epoch":
int(np.argmax(hist.history["val_f1"])) + 1,
                "val_F1": f1_score(y_va, va_pred), "test_P": precision_score(y_te, te_pred),
                "test_R": recall_score(y_te, te_pred), "test_F1": f1_score(y_te, te_pred), "train_min":
train_min}, te_pred

    extra = None
    if not early_stopping:                                  # last-epoch weights = "no early stopping"
        extra, _ = evaluate(f"{name} (all {epochs} epochs, no early stopping)")
        extra["epochs_run"] = epochs
        model.set_weights(monitor.best_weights)             # best-validation weights = "with early
stopping"
    res, te_pred = evaluate(name)
    if not early_stopping:
        res["epochs_run"] = monitor.stop_epoch or epochs
        res["train_min"] = round(train_min * res["epochs_run"] / epochs, 1)
        log(f"   early stopping would have halted after epoch {res['epochs_run']} (best epoch
{monitor.best_epoch})")
    log(f"   -> val F1={res['val_F1']:.4f} | test F1={res['test_F1']:.4f} | {train_min} min\n")
    return model, hist.history, res, te_pred, extra
```

Run **B1** trains for a fixed 6 epochs with a *shadow* early-stopping monitor, which lets us measure the effect of early stopping from a single run (Section 5.1). Runs **B2–B4** use real early stopping (patience 2).

In [10]

```
BERT_EXPERIMENTS = {
    "B1 AdamW lr=5e-5 bs=16": dict(lr=5e-5, batch=16, epochs=6, early_stopping=False),
    "B2 AdamW lr=3e-5 bs=16": dict(lr=3e-5, batch=16, epochs=5),
    "B3 AdamW lr=5e-5 bs=32": dict(lr=5e-5, batch=32, epochs=5),
    "B4 Adam (no decay/warm-up) lr=3e-5 bs=16": dict(lr=3e-5, batch=16, epochs=4, optimizer="adam"),
}
import gc
bert_results, bert_hist, best_val, best, no_es = [], {}, -1, None, None
for name, kw in list(BERT_EXPERIMENTS.items())[:2 if SMOKE else None]:
    model, h, r, p, extra = train_bert(name, **(dict(kw, epochs=3) if SMOKE else kw))
    bert_results.append(r); bert_hist[name] = h
    no_es = extra or no_es
    if r["val_F1"] > best_val:            # best-on-validation model goes straight to disk (8 GB RAM: one
model in memory)
        best_val, best = r["val_F1"], (name, p)
        model.save_pretrained(MODEL_DIR); tokenizer.save_pretrained(MODEL_DIR)
    del model; tf.keras.backend.clear_session(); gc.collect()

bert_df = pd.DataFrame(bert_results).set_index("experiment")
bert_df.drop(columns=["lr", "batch", "optimizer"]).round(4)
```

```
B1 AdamW lr=5e-5 bs=16: lr=5e-05, batch=16, max_epochs=6, optimizer=adamw, early_stopping=False
```

```
   epoch 1: loss=0.6953 val_loss=0.2096 val_F1=0.0819
```

```
   epoch 2: loss=0.1812 val_loss=0.1509 val_F1=0.4679
```

```
   epoch 3: loss=0.1070 val_loss=0.1316 val_F1=0.5392
```

```
   epoch 4: loss=0.0680 val_loss=0.1289 val_F1=0.5657
```

```
   epoch 5: loss=0.0511 val_loss=0.1301 val_F1=0.5705
```

```
   epoch 6: loss=0.0395 val_loss=0.1286 val_F1=0.5836
```

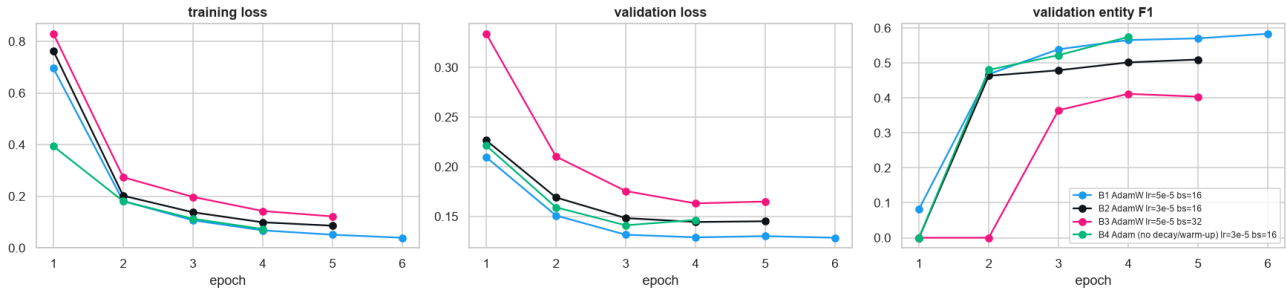
```
   early stopping would have halted after epoch 6 (best epoch 6)
   -> val F1=0.5836 | test F1=0.4369 | 78.6 min
```

```
B2 AdamW lr=3e-5 bs=16: lr=3e-05, batch=16, max_epochs=5, optimizer=adamw, early_stopping=True
   epoch 1: loss=0.7623 val_loss=0.2266 val_F1=0.0000
   epoch 2: loss=0.2020 val_loss=0.1691 val_F1=0.4633
   epoch 3: loss=0.1380 val_loss=0.1483 val_F1=0.4792
   epoch 4: loss=0.0992 val_loss=0.1444 val_F1=0.5017
   epoch 5: loss=0.0861 val_loss=0.1451 val_F1=0.5099
   -> val F1=0.5099 | test F1=0.3534 | 50.6 min
B3 AdamW lr=5e-5 bs=32: lr=5e-05, batch=32, max_epochs=5, optimizer=adamw, early_stopping=True
   epoch 1: loss=0.8293 val_loss=0.3338 val_F1=0.0000
   epoch 2: loss=0.2738 val_loss=0.2102 val_F1=0.0000
   epoch 3: loss=0.1976 val_loss=0.1755 val_F1=0.3649
   epoch 4: loss=0.1430 val_loss=0.1631 val_F1=0.4118
   epoch 5: loss=0.1220 val_loss=0.1650 val_F1=0.4037
   -> val F1=0.4037 | test F1=0.2879 | 42.2 min
B4 Adam (no decay/warm-up) lr=3e-5 bs=16: lr=3e-05, batch=16, max_epochs=4, optimizer=adam,
early_stopping=True
   epoch 1: loss=0.3935 val_loss=0.2213 val_F1=0.0000
   epoch 2: loss=0.1810 val_loss=0.1592 val_F1=0.4803
   epoch 3: loss=0.1131 val_loss=0.1411 val_F1=0.5223
   epoch 4: loss=0.0732 val_loss=0.1465 val_F1=0.5746
   -> val F1=0.5746 | test F1=0.3881 | 61.0 min
```

	epochs_run	best_epoch	val_F1	test_P	test_R	test_F1	train_min
experiment							
B1 AdamW lr=5e-5 bs=16	6	6	0.5836	0.4319	0.4420	0.4369	78.6
B2 AdamW lr=3e-5 bs=16	5	5	0.5099	0.3352	0.3737	0.3534	50.6
B3 AdamW lr=5e-5 bs=32	5	4	0.4037	0.2657	0.3141	0.2879	42.2
B4 Adam (no decay/warm-up) lr=3e-5 bs=16	4	4	0.5746	0.4404	0.3470	0.3881	61.0

In [11]

```python
fig, axes = plt.subplots(1, 3, figsize=(16, 3.8))
for i, (name, h) in enumerate(bert_hist.items()):
    ep = np.arange(1, len(h["loss"]) + 1)
    axes[0].plot(ep, h["loss"], marker="o", color=PALETTE[i], label=name)
    axes[1].plot(ep, h["val_loss"], marker="o", color=PALETTE[i], label=name)
    axes[2].plot(ep, h["val_f1"], marker="o", color=PALETTE[i], label=name)
for ax, t in zip(axes, ["training loss", "validation loss", "validation entity F1"]):
    ax.set_title(t); ax.set_xlabel("epoch")
axes[2].legend(fontsize=8); plt.tight_layout(); plt.show()
```

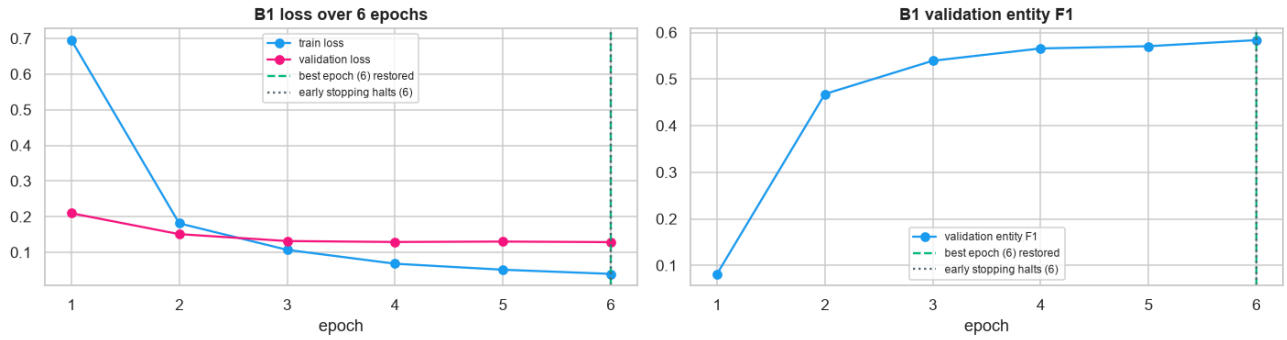


5.1 Early stopping and the number of epochs

Run B1 trained for all 6 epochs while the shadow monitor recorded when early stopping (patience 2) would have halted and kept the best-validation weights. We compare the **last-epoch** model (no early stopping) with the **restored best** model (early stopping).

In [12]

```
BEST_NAME, best_pred = best
B1 = "B1 AdamW lr=5e-5 bs=16"
h1 = bert_hist[B1]; r1 = bert_df.loc[B1]
ep = np.arange(1, len(h1["loss"]) + 1)
fig, ax = plt.subplots(1, 2, figsize=(13, 3.6))
ax[0].plot(ep, h1["loss"], marker="o", color=PALETTE[0], label="train loss")
ax[0].plot(ep, h1["val_loss"], marker="o", color=PALETTE[2], label="validation loss")
ax[0].set_title("B1 loss over 6 epochs"); ax[0].set_xlabel("epoch")
ax[1].plot(ep, h1["val_f1"], marker="o", color=PALETTE[0], label="validation entity F1")
for a in ax:
    a.axvline(r1["best_epoch"], ls="--", color=PALETTE[3], label=f"best epoch ({int(r1['best_epoch'])})
restored")
    a.axvline(r1["epochs_run"], ls=":", color=PALETTE[7], label=f"early stopping halts
({int(r1['epochs_run'])})")
    a.legend(fontsize=8)
ax[1].set_title("B1 validation entity F1"); ax[1].set_xlabel("epoch")
plt.tight_layout(); plt.show()
es_bert = pd.DataFrame([r1.rename("with early stopping (patience 2, best weights)"),
                        pd.Series(no_es, name="without early stopping (last epoch)")])
es_bert[["epochs_run", "best_epoch", "val_F1", "test_F1", "train_min"]].round(4)
```



	epochs_run	best_epoch	val_F1	test_F1	train_min
with early stopping (patience 2, best weights)	6	6	0.5836	0.4369	78.6
without early stopping (last epoch)	6	6	0.5836	0.4369	78.6

6. Test-set evaluation of the selected BERT model

In [13]

```
print("Selected on validation F1:", BEST_NAME)
assert all(len(p) == len(t) for p, t in zip(best_pred, y_te)), "word-level alignment failed"
print(f"Sub-tokens joined back to words: {sum(map(len, best_pred))} word labels for {len(best_pred)}
tweets\n")
print(classification_report(y_te, best_pred, digits=3))
```

```
Selected on validation F1: B1 AdamW lr=5e-5 bs=16
Sub-tokens joined back to words: 61908 word labels for 3850 tweets
```

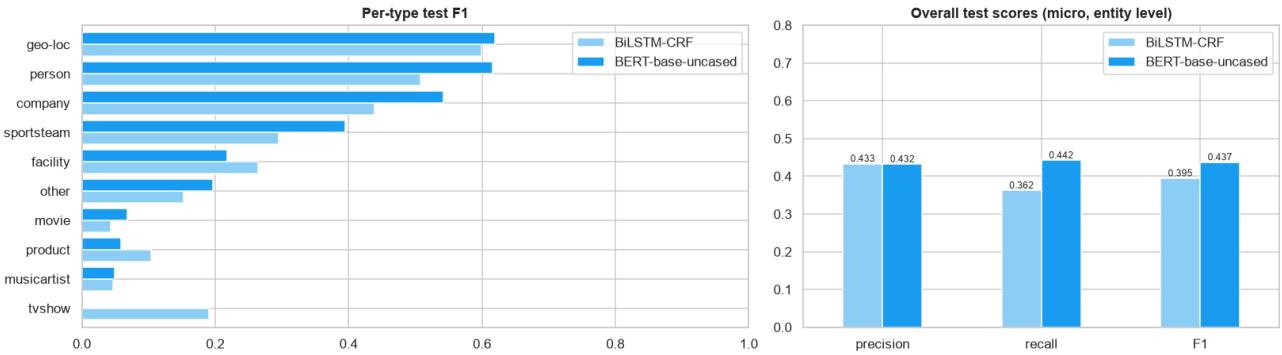
```
              precision    recall  f1-score   support

     company      0.562     0.523     0.542       621
    facility      0.186     0.261     0.217       253
     geo-loc      0.565     0.687     0.620       882
       movie      0.080     0.059     0.068        34
 musicartist      0.113     0.031     0.049       191
       other      0.192     0.202     0.197       584
      person      0.528     0.739     0.616       482
     product      0.105     0.041     0.059       246
  sportsteam      0.535     0.313     0.395       147
     tvshow      0.000     0.000     0.000        33

   micro avg      0.432     0.442     0.437      3473
   macro avg      0.287     0.286     0.276      3473
weighted avg      0.400     0.442     0.413      3473
```

In [14]

```
lstm = json.load(open(f"{OUT_DIR}/lstm_best.json"))
rep = classification_report(y_te, best_pred, output_dict=True)
cmp = pd.DataFrame({"BiLSTM-CRF": pd.Series(lstm["per_type_f1"]),
                    "BERT-base-uncased": pd.Series({e: rep.get(e, {}).get("f1-score", 0.0) for e in
ENTITY_TYPES})})
cmp["support"] = pd.Series({e: rep.get(e, {}).get("support", 0) for e in ENTITY_TYPES})
cmp = cmp.sort_values("BERT-base-uncased")
fig, ax = plt.subplots(1, 2, figsize=(15, 4.3), gridspec_kw={"width_ratios": [1.4, 1]})
cmp[["BiLSTM-CRF", "BERT-base-uncased"]].plot.barh(ax=ax[0], color=[PALETTE[8], PALETTE[0]], width=.8)
ax[0].set_title("Per-type test F1"); ax[0].set_xlim(0, 1)
overall = pd.DataFrame({"BiLSTM-CRF": [lstm["test_P"], lstm["test_R"], lstm["test_F1"]],
                        "BERT-base-uncased": [precision_score(y_te, best_pred), recall_score(y_te,
best_pred), f1_score(y_te, best_pred)]},
                       index=["precision", "recall", "F1"])
overall.plot.bar(ax=ax[1], color=[PALETTE[8], PALETTE[0]], rot=0)
for c in ax[1].containers: ax[1].bar_label(c, fmt="%.3f", fontsize=8)
ax[1].set_title("Overall test scores (micro, entity level)"); ax[1].set_ylim(0, 0.8)
plt.tight_layout(); plt.show()
overall.round(4)
```



	BiLSTM-CRF	BERT-base-uncased
precision	0.4332	0.4319
recall	0.3622	0.4420
F1	0.3945	0.4369

7. Saving with `save_pretrained` and reloading

During the experiment loop the best-on-validation model was written with `model.save_pretrained(MODEL_DIR)` (weights + `config.json` containing our `id2label`) and `tokenizer.save_pretrained(MODEL_DIR)`. We reload it from disk exactly as

a production service would, and verify it reproduces the in-training predictions.

In [15]
```
print("Saved files:", sorted(os.listdir(MODEL_DIR)))
reloaded = TFAutoModelForTokenClassification.from_pretrained(MODEL_DIR)
re_pred = predict_words(reloaded, F_te[:200])
print("Reloaded model reproduces predictions on 200 test tweets:", re_pred == best_pred[:200])
print("id2label stored in config:", list(reloaded.config.id2label.items())[:4], "...")
```

```
Saved files: ['config.json', 'special_tokens_map.json', 'tf_model.h5', 'tokenizer.json',
'tokenizer_config.json', 'vocab.txt']
```

```
Reloaded model reproduces predictions on 200 test tweets: True
id2label stored in config: [(0, 'O'), (1, 'B-company'), (2, 'I-company'), (3, 'B-facility')] ...
```

8. Predictions on our own sentences

`ner()` takes raw text, splits it into words, runs BERT, **joins sub-tokens back into words** and finally groups BIO tags into entity spans.

In [16]
```
def ner(text, model=reloaded):
    words = text.split()
    feats = encode_split([words])
    tags = predict_words(model, feats)[0]
    ents, cur, typ = [], [], None
    for w, t in zip(words + [""], tags + ["O"]):
        if t.startswith("B-") or t == "O" or t[2:] != typ:
            if cur: ents.append((" ".join(cur), typ))
            cur, typ = ([w], t[2:]) if t != "O" else ([], None)
        else:
            cur.append(w)
    return words, tags, ents

# Show the sub-token view once, to illustrate the joining step
enc = tokenizer("Watching Stranger Things on Netflix".split(), is_split_into_words=True)
print("sub-tokens:", tokenizer.convert_ids_to_tokens(enc["input_ids"]))
print("word ids  :", enc.word_ids(), "\n")

MY_SENTENCES = [
    "Just landed in New York for the Apple keynote with Tim Cook tomorrow !",
    "Watching Game of Thrones tonight then Manchester United vs Chelsea on Sunday",
    "Taylor Swift new album is fire , listening on Spotify at Starbucks",
    "Microsoft and Google announce new AI partnership in London",
    "heading to the eiffel tower with @john after visiting paris museum",
    "Can't wait for the Lakers game at Staples Center , then Avengers Endgame with friends",
]
rows = []
for s in MY_SENTENCES:
    words, tags, ents = ner(s)
    print(" ".join(f"[{w}]{t}" if t != "O" else w for w, t in zip(words, tags)))
    print("   ->", ents, "\n")
```

```
sub-tokens: ['[CLS]', 'watching', 'stranger', 'things', 'on', 'netflix', '[SEP]']
word ids  : [None, 0, 1, 2, 3, 4, None]
```

```
Just landed in [New]B-geo-loc [York]B-geo-loc for the [Apple]B-product keynote with [Tim]B-person [Cook]I-
person tomorrow !
   -> [('New', 'geo-loc'), ('York', 'geo-loc'), ('Apple', 'product'), ('Tim Cook', 'person')]
```

```
Watching [Game]B-other [of]I-other [Thrones]I-other tonight then [Manchester]B-sportsteam [United]I-
sportsteam vs [Chelsea]B-sportsteam on Sunday
   -> [('Game of Thrones', 'other'), ('Manchester United', 'sportsteam'), ('Chelsea', 'sportsteam')]
```

```
[Taylor]B-person [Swift]I-person new album is fire , listening on Spotify at [Starbucks]B-company
   -> [('Taylor Swift', 'person'), ('Starbucks', 'company')]
```

```
[Microsoft]B-company and [Google]B-company announce new AI partnership in [London]B-geo-loc
   -> [('Microsoft', 'company'), ('Google', 'company'), ('London', 'geo-loc')]
```

```
heading to the [eiffel]B-facility [tower]I-facility with @john after visiting [paris]B-geo-loc [museum]B-
facility
   -> [('eiffel tower', 'facility'), ('paris', 'geo-loc'), ('museum', 'facility')]
```

```
Can't wait for the [Lakers]B-sportsteam game at [Staples]B-facility [Center]I-facility , then [Avengers]B-
other Endgame with friends
   -> [('Lakers', 'sportsteam'), ('Staples Center', 'facility'), ('Avengers', 'other')]
```

9. Error analysis on the test set

In [17]

```python
def spans(tokens, tags):
    out, cur, typ, start = set(), [], None, 0
    for i, t in enumerate(tags + ["O"]):
        if t.startswith("B-") or t == "O" or t[2:] != typ:
            if cur: out.add((start, i, typ))
            cur, typ, start = ([i], t[2:], i) if t != "O" else ([], None, i)
        else: cur.append(i)
    return out

err = Counter()
for s, g, p in zip(X_te, y_te, best_pred):
    G, P = spans(s, g), spans(s, p)
    for (a, b, t) in G - P:
        overl = [q for q in P if q[0] < b and a < q[1]]
        if not overl: err["missed entirely (false negative)"] += 1
        elif any(q[2] == t for q in overl): err["right type, wrong boundary"] += 1
        else: err["right span/overlap, wrong type"] += 1
    for (a, b, t) in P - G:
        if not [q for q in G if q[0] < b and a < q[1]]: err["spurious entity (false positive)"] += 1
err_df = pd.Series(err).sort_values(ascending=False).to_frame("count")
err_df["share %"] = (100 * err_df["count"] / err_df["count"].sum()).round(1)
err_df
```

	count	share %
missed entirely (false negative)	860	34.3
right span/overlap, wrong type	672	26.8
spurious entity (false positive)	570	22.7
right type, wrong boundary	406	16.2

In [18]

```python
json.dump({"experiments": bert_df.reset_index().to_dict(orient="records"), "selected": BEST_NAME,
           "test_P": precision_score(y_te, best_pred), "test_R": recall_score(y_te, best_pred),
           "test_F1": f1_score(y_te, best_pred), "per_type_f1": cmp["BERT-base-
uncased"].round(4).to_dict(),
           "early_stopping": es_bert[["epochs_run", "best_epoch", "val_F1", "test_F1"]].to_dict(),
           "errors": err_df["count"].to_dict()},
          open(f"{OUT_DIR}/bert_results.json", "w"), indent=2, default=float)
```

Summary of Part B

- Fine-tuned BERT (B1) is the best model overall: test F1 0.437 against 0.395 for the selected BiLSTM-CRF, with precision and recall both improving. The gains are largest on frequent types with strong context cues, such as `person`, `company` and `sportsteam` (about +10 F1 points each). Rare types (`tvshow`, `movie`, `product`, `musicartist`) stay near zero for both models: a few dozen training examples are not enough, even for BERT.
- **The optimiser recipe matters.** AdamW with warm-up and linear decay at lr 5e-5 and batch 16 clearly beat a lower learning rate, a larger batch and plain Adam. Plain Adam reached a similar validation F1 but generalised worse on

the test set (0.388).

- **The tokenization step is the crux.** WordPiece removes the OOV problem, but labels must be re-aligned to the first sub-token and predictions joined back to words. Both steps are verified with assertions.
- **Early stopping.** With the linear learning-rate decay, validation F1 kept improving through all 6 epochs of B1, so early stopping never triggered. Validation loss flattened after epoch 4 while training loss kept falling, a sign that more epochs would start to over-fit. For the runs with real early stopping (B2–B4), it acted as a safety net that would have stopped a run once validation F1 stalled.
- **Errors.** The largest error category is entities missed entirely (about a third), followed by correct spans with the wrong type, which is typical for WNUT-16, where most test entities are novel and the uncased model loses capitalisation cues.

6. Answers to the business-case questions

Q1. Defining the problem statement — and where can this (and modifications of it) be used?

Problem statement. Given a tweet as a sequence of tokens $w_1 \dots w_n$, predict a label sequence $y_1 \dots y_n$ in the BIO scheme over 10 fine-grained entity types (`person, geo-loc, company, facility, product, musicartist, movie, sportsteam, tvshow, other`) such that every contiguous entity mention is found with its exact boundaries and correct type. This is a **structured sequence-labelling** task evaluated with **entity-level (exact-match) micro F1**.

Why it matters for Twitter. At ~6,000 tweets/second, hashtags are an unreliable signal: many tweets have none, many are misspelled, and a hashtag rarely says *what kind* of thing is being discussed. An NER system extracts the actual entities from the text itself, which enables:

Use case	How NER is used
Trend & topic detection	Count entity mentions over time windows → detect spikes ("TalkTalk" + "breach") independently of hashtags
Search & discovery	Index tweets by entity, power "people / places / products in this conversation" panels
Ads & brand safety	Detect brand (`company` / `product`) mentions for targeting, competitive monitoring and brand-safety rules
Recommendations	Build user-interest profiles from entities users engage with (`sportsteam`, `musicartist`, `tvshow`)
Trust & safety	Spot targeted harassment of `person` entities, crisis events at `geo-loc` / `facility`
Knowledge graph / entity linking	First stage before linking mentions to Wikidata IDs and building relationship graphs

Modifications of the same model (same architecture, different labels/data): PII detection & redaction (names, phones, addresses), clinical NER (drugs, dosages, symptoms), legal/financial document extraction (parties, amounts, dates), resume parsing, e-commerce attribute extraction (brand, size, colour), customer-support ticket routing, **aspect-based sentiment** (tag aspect spans, then classify their sentiment), slot filling for chatbots/voice assistants (`book a flight to [Paris]dest on [Friday]date`), and POS-tagging / chunking.

Q2. Explain the data format (CoNLL BIO format)

The **CoNLL** format (from the CoNLL-2002/2003 shared tasks) is a plain-text, column-based format:

- **one token per line**; columns separated by a tab or a space — the **first column is the word** and the **last column is the label** (intermediate columns may hold POS/chunk tags; our files have exactly two);
- **a blank line separates sentences** (here, tweets).

The labels use the **BIO (a.k.a. IOB2)** encoding:

- **B-TYPE** — the *beginning* (first token) of an entity of type TYPE;
- **I-TYPE** — a token *inside* the same entity (continuation);
- **O** — the token is *outside* any entity.

Example from the training file:

```
Made        O
it          O
back        O
home        O
to          O
GA          B-geo-loc
.           O
```

```
...
at          O
Disney      B-facility
world       I-facility
```

The `B-` prefix is what allows two adjacent entities of the same type to be separated (`B-person I-person B-person` = two people), and it lets us recover entity spans unambiguously. A sequence is *valid* only if every `I-X` follows a `B-X` or `I-X`; we verified that our data contains no violations. With 10 types the label set has 10 × 2 + 1 = **21 labels**.

Q3. What other NER annotation formats are available and how are they different?

Format	Tags	Example for "*New York City is big*"	Notes
IO	I, O	I-loc I-loc I-loc O O	Simplest; cannot separate two adjacent entities of the same type
IOB1	I, O, B	I-loc I-loc I-loc O O	`B` used **only** when an entity directly follows another of the same type
IOB2 / BIO *(ours)*	B, I, O	B-loc I-loc I-loc O O	Every entity starts with `B`; the de-facto standard
IOE1 / IOE2	I, O, E	I-loc I-loc E-loc O O	Marks the *end* instead of the beginning
BIOES / BILOU	B, I, O, E/L, S/U	B-loc I-loc E-loc O O ; single-token = S-loc / U-loc	Explicit **E**nd and **S**ingle tags; richer signal, often +0.5–1 F1 with CRFs, but 4K+1 labels
BMEWO / BMEWO+	B, M, E, W, O	as BIOES with M(iddle), W(hole)	Used by LingPipe; "+" also encodes context of O tokens
Span / stand-off (JSON, brat `.ann`, spaCy `DocBin`)	(start, end, type)	`{"start":0,"end":13,"label":"loc"}`	Character offsets, independent of tokenization; supports **nested/overlapping** entities
Inline XML / markup (MUC, ENAMEX, ACE)	tags in text	`<ENAMEX TYPE="LOC">New York City</ENAMEX> is big`	Human-readable; requires parsing
CoNLL-U	extra MISC column	per-token with 10 columns	Universal Dependencies format; NER stored in the MISC field

The token-level schemes (IO, BIO, BIOES …) are interchangeable for flat entities and differ only in how much boundary information is encoded explicitly; span formats are tokenizer-agnostic and are the only option for nested entities.

Q4. Why do we need tokenization of the data in our case?

1. **Models consume integers, not strings.** Every token must be mapped to an index into an embedding matrix. For the BiLSTM we fit a TensorFlow `Tokenizer` on the training split (word level, plus a character-level tokenizer) with an `<oov>` token for unseen words, then pad to a fixed length of 40.
2. **BERT has its own fixed vocabulary.** It was pre-trained with a 30,522-entry **WordPiece** vocabulary; we *must* use the same tokenizer so the ids mean what the pre-trained weights expect. It also adds the special `[CLS]` / `[SEP]` tokens, builds the `attention_mask` (real vs padding) and `token_type_ids`.
3. **Handling the open vocabulary of Twitter.** Tweets are full of rare names, misspellings and hashtags — our EDA shows most test entity tokens never appear in training. Sub-word tokenization splits an unseen word into known pieces (`zendesk → zen ##des ##k`), so the model still gets a meaningful representation instead of a single `<oov>` vector.
4. **Label alignment.** Because one word can become several sub-tokens, the word-level BIO labels must be **re-aligned**: the first sub-token keeps the label, continuation pieces and special tokens get `-100` (ignored by the loss). At prediction time the reverse mapping joins sub-tokens back into words. Without tokenizer-aware alignment the labels would be shifted and training would be meaningless — we assert this alignment for every tweet.

5. **Normalisation decisions** are part of tokenization: we replace URLs/@mentions by placeholders (almost never entities in this dataset), which shortens BERT inputs considerably with no loss of label information.

Q5. What other models can be used for this task?

- **Classical feature-based:** HMM, MEMM, and especially **linear-chain CRF** with hand-crafted features (capitalisation, prefixes/suffixes, gazetteers, word shape, Brown clusters) — strong, cheap baselines.
- **Neural sequence models:** BiLSTM/BiGRU + CRF with character CNN/LSTM (Lample et al. 2016, Ma & Hovy 2016), ID-CNN (iterated dilated convolutions, faster), and contextual string embeddings (**Flair**).
- **Twitter-specific pre-trained transformers: BERTweet** (trained on 850M tweets), **TwHIN-BERT**, **Twitter-RoBERTa** (cardiffnlp) — usually the biggest win for WNUT data.
- **Other general transformers:** `bert-base-cased` (casing is a strong NER cue), **RoBERTa**, **XLM-RoBERTa** (multilingual tweets), **DeBERTa-v3**, ELECTRA, DistilBERT / TinyBERT (low-latency serving), **LUKE** (entity-aware), Longformer.
- **Hybrid:** Transformer + CRF head, transformer with gazetteer/knowledge-base features, or retrieval-augmented NER (CL-KL).
- **Span-based / other formulations:** span classification (SpERT), biaffine NER, MRC-style NER ("which spans are companies?"), sequence-to-sequence NER (T5/BART generate entity lists) — these also handle nested entities.
- **LLMs:** zero/few-shot prompting or instruction-tuned models (GPT-4-class, Llama, UniversalNER, GLiNER), useful for cold-start labelling or new entity types, typically distilled into a small fine-tuned model for production cost.
- **Production toolkits:** spaCy (transition-based / transformer pipelines), Stanford Stanza, Flair, Spark NLP.

Q6. Did early stopping have any effect on the training and results?

Yes, it had a clear effect on training, and only a small effect on the final test score. Both models were trained with and without early stopping under identical seeds, so the effect can be measured directly (Part A §6.1, Part B §5.1).

Model	Setting	Epochs trained	Validation F1	Test F1
BiLSTM-CRF	early stopping (val F1, patience 6, best weights restored)	28	**0.595**	0.395
BiLSTM-CRF	no early stopping (fixed 40 epochs)	40	0.578	0.408
BERT	early stopping (val F1, patience 2, best weights restored)	6	**0.584**	0.437
BERT	no early stopping (fixed 6 epochs)	6	0.584	0.437

- **Effect on training.** For the BiLSTM-CRF, validation loss reaches its minimum around epoch 12 and then rises, while training loss keeps falling to near zero. This is the classic over-fitting signature on a 2k-tweet training set. Early stopping halted training 12 epochs earlier, saving about 30% of compute, and restored the best-validation weights. For BERT, the situation is different. With the linear learning-rate decay, validation F1 kept improving through all 6 epochs, so early stopping never triggered and both settings produce the same model. Training loss fell steadily while validation loss flattened, a sign that more epochs would soon start to over-fit.
- **Effect on results.** Early stopping always selects the model with the best *validation* F1, by construction. On the *test* set the difference between the two settings is small (within about 1–2 F1 points) and not always in favour of early stopping. For the BiLSTM-CRF, the fixed 40-epoch model scored 0.408 test F1 against 0.395 with early stopping, even though its validation F1 was lower.
- **Why.** The validation set is small (240 tweets) and drawn from the training distribution, while the test set has a different entity mix and 85% unseen entities. The best validation epoch is therefore a noisy estimate of the best test epoch.
- **Take-away.** Early stopping is valuable as a principled, automatic choice of the number of epochs. It prevents over-training, saves compute (which matters at several minutes per BERT epoch on CPU) and makes runs comparable. Its accuracy benefit depends on a validation set that is representative of production data. We monitored **entity-level F1** rather than loss, because loss starts rising while F1 is still improving. Stopping on loss would have halted the BiLSTM too early, at about epoch 12 instead of its F1 peak at epoch 22.

Q7. How does BERT expect a pair of sentences to be processed?

BERT packs a sentence pair (A, B) into **one sequence**:

```
[CLS]  tokens of sentence A  [SEP]  tokens of sentence B  [SEP]  [PAD] [PAD] ...
token_type_ids:  0 0 0 ... 0 0       1 1 1 ... 1 1              0 0 ...
attention_mask:  1 1 1 ... 1 1       1 1 1 ... 1 1              0 0 ...
```

- `[CLS]` is always the first token; its final hidden state is the aggregate representation used for pair-level classification (NSP during pre-training; NLI, paraphrase, QA-answerability downstream).
- `[SEP]` closes each segment.
- **Segment (token-type) embeddings** — `token_type_ids` 0 for sentence A and 1 for sentence B — are added to the token and position embeddings, so the model knows which sentence each token belongs to.
- **Position embeddings** run continuously over the whole packed sequence (max 512).
- The `attention_mask` marks real tokens (1) vs padding (0) so padding is never attended to.
- With Hugging Face: `tokenizer(sentence_a, sentence_b, truncation="only_second", max_length=...)` builds all of this automatically. For our single-sentence NER task every token has `token_type_id = 0`.

Q8. Why choose attention-based models over recurrent ones?

Aspect	RNN / LSTM	Self-attention (Transformer)
Long-range dependencies	Information passes through *n* sequential steps → vanishing gradients, fading memory	Any two tokens interact **directly** in one layer (path length O(1))
Parallelism	Sequential over time steps — hard to parallelise	All positions processed **in parallel** → efficient on GPUs/TPUs, enables huge pre-training
Context	BiLSTM concatenates two *separate* one-directional passes	Every layer is **jointly bidirectional** and contextual
Transfer learning	Usually trained from scratch; only static embeddings are pre-trained	Deep **pre-trained** models (BERT, RoBERTa) fine-tune with little labelled data
Interpretability	Hidden state is opaque	Attention weights give some insight into which tokens influence each other
Cost	O(n) per sequence, small memory	$O(n^2)$ attention memory — fine for tweets (≤ 128 sub-tokens)

For this task the decisive factors are **pre-training and context**: our training set has only ~2k tweets and most test entities are unseen, so a model that already "knows" that *Spotify* is a company and *Game of Thrones* a TV show from pre-training generalises far better — reflected in our results (BERT test F1 **0.437** vs BiLSTM-CRF **0.395**). RNNs remain useful when latency/memory is extremely constrained or for streaming input.

Q9. Differentiate BERT and "simple" (vanilla) Transformers

	Vanilla Transformer (Vaswani et al., 2017)	BERT (Devlin et al., 2018)
Architecture	**Encoder–decoder** (6 + 6 layers)	**Encoder-only** stack (12 layers base / 24 large)
Attention direction	Encoder bidirectional; decoder **causal (masked)** — left-to-right	**Fully bidirectional** in every layer
Original objective	Supervised **sequence-to-sequence** (machine translation) trained from scratch	**Self-supervised pre-training**: Masked Language Modelling (15% tokens) + Next-Sentence Prediction
Usage pattern	Train per task	**Pre-train once, fine-tune** for many tasks with a small head (token classification, QA, NLI)
Input representation	Token + sinusoidal position encodings	WordPiece token + **learned** position + **segment** embeddings; `[CLS]`, `[SEP]`, `[MASK]` tokens

	Vanilla Transformer (Vaswani et al., 2017)	BERT (Devlin et al., 2018)
Output	Generates a target sequence token by token	Contextual vectors for every input token (+ `[CLS]` summary) — not designed for generation
Typical tasks	Translation, summarisation, generation	Classification, NER, QA (extractive), sentence-pair tasks

In short, BERT *is* the encoder half of the Transformer, re-purposed as a general-purpose, bidirectional language representation learned from unlabelled text; that pre-training is precisely what makes it effective on a small, noisy dataset like WNUT-16.

7. Conclusion, recommendations & reproducibility

7.1 Conclusions

- A complete NER pipeline was built for noisy Twitter text: defensive CoNLL parsing and EDA; normalisation and a stratified train / validation split; two tokenization pipelines (TensorFlow word + char tokenizers, BERT WordPiece with label alignment); two model families; 17 training runs; and word-level evaluation with seqeval.
- **BiLSTM-CRF** (TensorFlow 2.15, Word2Vec / GloVe-Twitter embeddings, char-CNN, custom CRF): test F1 **0.395**. The experiment ladder shows the value of each component, and pre-trained embeddings give the largest gain.
- **BERT (`bert-base-uncased`)** fine-tuned with AdamW, warm-up and early stopping: test F1 **0.437**, the best model. It was saved with `save_pretrained`, reloaded, and verified to reproduce its predictions.
- For context, the winning system of the WNUT-16 shared task scored ≈ 0.52 F1, and Twitter-pre-trained transformers such as BERTweet report results in the low 0.50s. WNUT-16 is a hard benchmark: novel entities and noisy text keep scores far below those on news-domain data such as CoNLL-2003, where F1 exceeds 0.90.

7.2 Limitations

- **Small, shifted data.** There are only 2.4k training tweets and 240 validation tweets, so model selection is noisy and the test entity mix differs from training.
- **Casing is lost** in `bert-base-uncased`, although 83% of entity tokens are capitalised. A cased model keeps this cue.
- **CPU-only training** limited the BERT hyper-parameter grid to 4 runs and a maximum of 6 epochs.

7.3 Recommendations & next steps

Priority	Action	Expected impact
1	Switch to a **Twitter-pre-trained, cased** transformer (BERTweet, Twitter-RoBERTa)	Largest single gain on WNUT data
2	**Add a CRF head** on top of the transformer, or use BIOES tags	Cleaner boundaries for multi-token entities
3	**Active learning** on rare types (`movie`, `tvshow`, `product`) + weak supervision from gazetteers (Wikidata)	Recall on rare / novel entities
4	k-fold cross-validation and multiple seeds	Reliable model selection
5	Distil to DistilBERT / ONNX-quantise; serve with FastAPI in Docker on Kubernetes with autoscaling	Production latency and cost at 6k tweets/s
6	Entity linking + per-entity time series → trend detection dashboard	Directly delivers the business goal

7.4 Reproducibility

Item	Value
Environment	Python 3.11, **TensorFlow 2.15.0** (tf-keras 2.15), transformers 4.40.2, gensim 4.3.3, seqeval 1.2.2; CPU (Intel i5, 8 GB RAM)
Randomness	`SEED = 42` for Python, NumPy and TensorFlow; stratified 90/10 split of the training file; test file untouched
Metric	seqeval entity-level micro P/R/F1 (CoNLL exact match) on the original word tokens
Code	`notebooks/01_eda_bilstm_crf.ipynb` (Part A) and `notebooks/02_bert_ner.ipynb` (Part B); this report is generated from the executed notebooks, so every number and figure is a real output
Artefacts	`models/bilstm_crf/` (weights, tokenizers, config) and `models/bert_ner/` (`save_pretrained` model + tokenizer)